

Semantic Segmentation

1. Для начала мы скачаем датасет: [ADDI project \(https://www.fc.up.pt/addi/ph2%20database.html\)](https://www.fc.up.pt/addi/ph2%20database.html).



1. Разархивируем .rar файл.
2. Обратите внимание, что папка PH2 Dataset images должна лежать там же где и ipynb notebook.

Это фотографии двух типов **поражений кожи**: меланома и родинки. В данном задании мы не будем заниматься их классификацией, а будем сегментировать их.

In [0]:

```
import numpy as np
import torch

def seed_init(SEED = 1500):
    np.random.seed(SEED)
    torch.manual_seed(SEED)
    torch.cuda.manual_seed(SEED)
    torch.backends.cudnn.deterministic = True

seed_init()
```

Структура датасета у нас следующая:

```
IMD_002/
  IMD002_Dermoscopic_Image/
    IMD002.bmp
  IMD002_lesion/
    IMD002_lesion.bmp
  IMD002_roi/
    ...
IMD_003/
  ...
  ...
```

Для загрузки датасета я предлагаю использовать skimage: `skimage.io.imread()` (<https://scikit-image.org/docs/dev/api/skimage.io.html>)

In [0]:

```
images = []
lesions = []
from skimage.io import imread
import os
root = '/content/drive/My Drive/Colab Notebooks/Segmentation/data/PH2Dataset/'

for root, dirs, files in os.walk(os.path.join(root, 'PH2 Dataset images')):
    if root.endswith('_Dermoscopic_Image'):
        images.append(imread(os.path.join(root, files[0])))
    if root.endswith('_lesion'):
        lesions.append(imread(os.path.join(root, files[0])))
```

Изображения имеют разные размеры. Давайте изменим их размер на 256×256 пикселей.

`skimage.transform.resize()` (<https://scikit-image.org/docs/dev/api/skimage.transform.html#skimage.transform.resize>) можно использовать для изменения размера изображений. Эта функция также автоматически нормализует изображения в диапазоне $[0,1]$.

In [0]:

```
from skimage.transform import resize
size = (256, 256)
X = [resize(x, size, mode='constant', anti_aliasing=True,) for x in images]
Y = [resize(y, size, mode='constant', anti_aliasing=False) > 0.5 for y in lesions]
Y_ = [resize(y, size, mode='constant', anti_aliasing=False) for y in lesions]
```

In [4]:

```
import numpy as np
X = np.array(X, np.float32)
Y = np.array(Y, np.float32)
Y_ = np.array(Y_, np.float32)
print(f'Loaded {len(X)} images')
```

Loaded 199 images

In [5]:

```
len(lesions)
```

Out[5]:

199

Чтобы убедиться, что все корректно, мы нарисуем несколько изображений

In [6]:


```
import matplotlib.pyplot as plt
from IPython.display import clear_output

plt.figure(figsize=(18, 6))
for i in range(6):
    plt.subplot(3, 6, i+1)
    plt.axis("off")
    plt.imshow(X[i])

    plt.subplot(3, 6, i+7)
    plt.axis("off")
    plt.imshow(Y[i])

    plt.subplot(3, 6, i+13)
    plt.axis("off")
    plt.imshow(Y_[i])

plt.show();
```



Разделим наши 200 картинок на 100/50/50 для валидации и теста

In [0]:

```
ix = np.random.choice(len(X), len(X), False)
tr, val, ts = np.split(ix, [100, 150])
```

In [8]:

```
print(len(tr), len(val), len(ts))
```

100 50 49

PyTorch DataLoader

In [0]:

```
from torch.utils.data import DataLoader
batch_size = 25
data_tr = DataLoader(list(zip(np.rollaxis(X[tr], 3, 1), Y[tr, np.newaxis])),
                      batch_size=batch_size, shuffle=True)
data_val = DataLoader(list(zip(np.rollaxis(X[val], 3, 1), Y[val, np.newaxis])),
                      batch_size=batch_size, shuffle=True)
data_ts = DataLoader(list(zip(np.rollaxis(X[ts], 3, 1), Y[ts, np.newaxis])),
                      batch_size=batch_size, shuffle=True)
```

In [11]:

```
import torch
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device
```

Out[11]:

device(type='cuda')

In [0]:

```
import copy
def save_weights(net, name):
    net_weights = copy.deepcopy(net.state_dict())
    torch.save(net_weights, name)
    !mv '/content/{name}' '/content/drive/My Drive/Colab Notebooks/Segmentation/weights'
```


In [0]:

```
def load_weights(*nets, paths):
    for net, path in zip(nets, paths):
        net.load_state_dict(torch.load(path))
```

Реализация различных архитектур:

Ваше задание будет состоять в том, чтобы написать несколько нейросетевых архитектур для решения задачи семантической сегментации. Сравнить их по качеству на тесте и испробовать различные лосс функции для них.

SegNet [2 балла]



- Badrinarayanan, V., Kendall, A., & Cipolla, R. (2015). [SegNet: A deep convolutional encoder-decoder architecture for image segmentation](https://arxiv.org/pdf/1511.00561.pdf) (<https://arxiv.org/pdf/1511.00561.pdf>)

In [0]:

```
import gc
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from time import time
from torchsummary import summary

from matplotlib import rcParams
rcParams['figure.figsize'] = (15,8)
```

In [0]:

```
def clear_memory():
    gc.collect()
    torch.cuda.empty_cache()
```

In [0]:

```
class SegNet(nn.Module):
    def __init__(self):
        super().__init__()

        self.enc_conv0 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
        )

        self.pool0 = nn.MaxPool2d(kernel_size=2, stride=2) # 256 -> 128
        self.enc_conv1 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
        )

        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2) # 128 -> 64
        self.enc_conv2 = nn.Sequential(
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.3),
            nn.Conv2d(128, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
        )

        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2) # 64 -> 32
        self.enc_conv3 = nn.Sequential(
            nn.Conv2d(128, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.35),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
        )

        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2) # 32 -> 16

        # bottleneck
        self.bottleneck_conv = nn.Sequential(
            nn.Conv2d(256, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.45),
            nn.Conv2d(512, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
            nn.Conv2d(512, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512).
```

```

nn.ReLU(inplace=True),
)

# decoder (upsampling)
self.upsample0 = nn.Upsample(scale_factor=2, mode="bilinear") # 16 -> 32
self.dec_conv0 = nn.Sequential(
    nn.Conv2d(512, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.35),
    nn.Conv2d(256, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
    nn.Conv2d(256, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
)
self.upsample1 = nn.Upsample(scale_factor=2, mode="bilinear") # 32 -> 64
self.dec_conv1 = nn.Sequential(
    nn.Conv2d(256, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(128, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
)
self.upsample2 = nn.Upsample(scale_factor=2, mode="bilinear") # 64 -> 128
self.dec_conv2 = nn.Sequential(
    nn.Conv2d(128, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),
    nn.Conv2d(64, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),
)
self.upsample3 = nn.Upsample(scale_factor=2, mode="bilinear") # 128 -> 256
self.dec_conv3 = nn.Sequential(
    nn.Conv2d(64, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.ReLU(inplace=True),
    nn.Conv2d(32, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.ReLU(inplace=True),
    nn.Conv2d(32, 1, kernel_size=3, padding=1),
    nn.BatchNorm2d(1)
)

def forward(self, x):
    # encoder
    e0 = self.pool0(self.enc_conv0(x))
    e1 = self.pool1(self.enc_conv1(e0))
    e2 = self.pool2(self.enc_conv2(e1))
    e3 = self.pool3(self.enc_conv3(e2))

    # bottleneck
    b = self.bottleneck_conv(e3)

    # decoder
    d0 = self.dec_conv0(self.upsample0(b))
    d1 = self.dec_conv1(self.upsample1(d0))
    d2 = self.dec_conv2(self.upsample2(d1))
    d3 = self.dec_conv3(self.upsample3(d2)) # no activation
    return d3

```

In [17]:

```
summary(SegNet().to(device), input_size=(3, 256, 256))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	9,248
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
MaxPool2d-7	[-1, 32, 128, 128]	0
Conv2d-8	[-1, 64, 128, 128]	18,496

BatchNorm2d-9	[-1, 64, 128, 128]	128
ReLU-10	[-1, 64, 128, 128]	0
Conv2d-11	[-1, 64, 128, 128]	36,928
BatchNorm2d-12	[-1, 64, 128, 128]	128
ReLU-13	[-1, 64, 128, 128]	0
MaxPool2d-14	[-1, 64, 64, 64]	0
Conv2d-15	[-1, 128, 64, 64]	73,856
BatchNorm2d-16	[-1, 128, 64, 64]	256
ReLU-17	[-1, 128, 64, 64]	0
Dropout2d-18	[-1, 128, 64, 64]	0
Conv2d-19	[-1, 128, 64, 64]	147,584
BatchNorm2d-20	[-1, 128, 64, 64]	256
ReLU-21	[-1, 128, 64, 64]	0
MaxPool2d-22	[-1, 128, 32, 32]	0
Conv2d-23	[-1, 256, 32, 32]	295,168
BatchNorm2d-24	[-1, 256, 32, 32]	512
ReLU-25	[-1, 256, 32, 32]	0
Dropout2d-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 256, 32, 32]	590,080
BatchNorm2d-28	[-1, 256, 32, 32]	512
ReLU-29	[-1, 256, 32, 32]	0
Conv2d-30	[-1, 256, 32, 32]	590,080
BatchNorm2d-31	[-1, 256, 32, 32]	512
ReLU-32	[-1, 256, 32, 32]	0
MaxPool2d-33	[-1, 256, 16, 16]	0
Conv2d-34	[-1, 512, 16, 16]	1,180,160
BatchNorm2d-35	[-1, 512, 16, 16]	1,024
ReLU-36	[-1, 512, 16, 16]	0
Dropout2d-37	[-1, 512, 16, 16]	0
Conv2d-38	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-39	[-1, 512, 16, 16]	1,024
ReLU-40	[-1, 512, 16, 16]	0
Conv2d-41	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-42	[-1, 512, 16, 16]	1,024
ReLU-43	[-1, 512, 16, 16]	0
Upsample-44	[-1, 512, 32, 32]	0
Conv2d-45	[-1, 256, 32, 32]	1,179,904
BatchNorm2d-46	[-1, 256, 32, 32]	512
ReLU-47	[-1, 256, 32, 32]	0
Dropout2d-48	[-1, 256, 32, 32]	0
Conv2d-49	[-1, 256, 32, 32]	590,080
BatchNorm2d-50	[-1, 256, 32, 32]	512
ReLU-51	[-1, 256, 32, 32]	0
Conv2d-52	[-1, 256, 32, 32]	590,080
BatchNorm2d-53	[-1, 256, 32, 32]	512
ReLU-54	[-1, 256, 32, 32]	0
Upsample-55	[-1, 256, 64, 64]	0
Conv2d-56	[-1, 128, 64, 64]	295,040
BatchNorm2d-57	[-1, 128, 64, 64]	256
ReLU-58	[-1, 128, 64, 64]	0
Dropout2d-59	[-1, 128, 64, 64]	0
Conv2d-60	[-1, 128, 64, 64]	147,584
BatchNorm2d-61	[-1, 128, 64, 64]	256
ReLU-62	[-1, 128, 64, 64]	0
Upsample-63	[-1, 128, 128, 128]	0
Conv2d-64	[-1, 64, 128, 128]	73,792
BatchNorm2d-65	[-1, 64, 128, 128]	128
ReLU-66	[-1, 64, 128, 128]	0
Conv2d-67	[-1, 64, 128, 128]	36,928
BatchNorm2d-68	[-1, 64, 128, 128]	128
ReLU-69	[-1, 64, 128, 128]	0
Upsample-70	[-1, 64, 256, 256]	0
Conv2d-71	[-1, 32, 256, 256]	18,464
BatchNorm2d-72	[-1, 32, 256, 256]	64
ReLU-73	[-1, 32, 256, 256]	0
Conv2d-74	[-1, 32, 256, 256]	9,248
BatchNorm2d-75	[-1, 32, 256, 256]	64
ReLU-76	[-1, 32, 256, 256]	0
Conv2d-77	[-1, 1, 256, 256]	289
BatchNorm2d-78	[-1, 1, 256, 256]	2

```

=====
Total params: 10,611,459
Trainable params: 10,611,459
Non-trainable params: 0

```

```

-----
Input size (MB): 0.75
Forward/backward pass size (MB): 462.50
Params size (MB): 40.48
Estimated Total Size (MB): 503.73
-----

```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

"See the documentation of nn.Upsample for details.".format(mode))

Метрика

В данном разделе предлагается использовать следующую метрику для оценки качества:

$$IoU = \frac{\text{target} \cap \text{prediction}}{\text{target} \cup \text{prediction}}$$

Пересечение ($A \cap B$) состоит из пикселей, найденных как в маске предсказания, так и в основной маске истины, тогда как объединение ($A \cup B$) просто состоит из всех пикселей, найденных либо в маске предсказания, либо в целевой маске.

To clarify this we can see on the segmentation:



And the intersection will be the following:



In [0]:

```
def iou_pytorch(outputs: torch.Tensor, labels: torch.Tensor):
    # You can comment out this line if you are passing tensors of equal shape
    # But if you are passing output from UNet or something it will most probably
    # be with the BATCH x 1 x H x W shape
    outputs = outputs.squeeze(1).byte() # BATCH x 1 x H x W => BATCH x H x W
    labels = labels.squeeze(1).byte()
    SMOOTH = 1e-8
    intersection = (outputs & labels).float().sum((1, 2)) # Will be zero if Truth=0 or Prediction=0
    union = (outputs | labels).float().sum((1, 2)) # Will be zero if both are 0

    iou = (intersection + SMOOTH) / (union + SMOOTH) # We smooth our division to avoid 0/0

    thresholded = torch.clamp(20 * (iou - 0.5), 0, 10).ceil() / 10 # This is equal to comparing with thresholds

    return thresholded
```

функция лосса [1 балл]

Теперь не менее важным, чем построение архитектуры, является определение **оптимизатора** и **функции потерь**.

Функция потерь - это то, что мы пытаемся минимизировать. Многие из них могут быть использованы для задачи бинарной семантической сегментации.

Популярным методом для бинарной сегментации является *бинарная кросс-энтропия*, которая задается следующим образом:

$$\mathcal{L}_{\text{BCE}}(y, \hat{y}) = -\sum_i \left[y_i \log(\sigma(\hat{y}_i)) + (1 - y_i) \log(1 - \sigma(\hat{y}_i)) \right]$$

где y это таргет желаемого результата и \hat{y} является выходом модели. σ - это [логистическая функция](https://en.wikipedia.org/wiki/Sigmoid_function) (https://en.wikipedia.org/wiki/Sigmoid_function), который преобразует действительное число \mathbb{R} в вероятность $[0,1]$.

Однако эта потеря страдает от проблем численной нестабильности. Самое главное, что $\lim_{x \rightarrow 0} \log(x) = -\infty$ приводит к неустойчивости в процессе оптимизации. Рекомендуется посмотреть следующее [упрощение](https://www.tensorflow.org/api_docs/python/tf/nn/softmax_cross_entropy_with_logits) (https://www.tensorflow.org/api_docs/python/tf/nn/softmax_cross_entropy_with_logits) в Тарая функция эквивалентна и не так подвержена численной неустойчивости.

$$\mathcal{L}_{\text{BCE}} = \hat{y} - y + \log(1 + \exp(-\hat{y}))$$

In [0]:

```
def bce_loss(y_pred, y_real, reduce=True):
    # TODO
    # please don't use nn.BCELoss. write it from scratch
    y_pred = torch.sigmoid(y_pred)
    preout = y_pred.sub(y_real.mul(y_pred))
    out = preout.add((1 + (-y_pred).exp()).log())
    return out.mean() if reduce else out
```

Тренировка [1 балл]

Мы определим цикл обучения в функции, чтобы мы могли повторно использовать его.

In [0]:

```
def eval_epoch(model, val_loader, criterion, metric, threshold):
    running_loss = .0
    running_score = .0
    processed_size = len(val_loader)

    model = model.to(device)
    model.eval()
    for inputs, real_masks in val_loader:
        inputs = inputs.to(device)
        real_masks = real_masks.to(device)

        with torch.no_grad():
            outputs = model(inputs)
            loss = criterion(outputs, real_masks)
            pred_masks = torch.sigmoid(outputs) > threshold

        running_score += metric(pred_masks, real_masks).mean().item()
        running_loss += loss.item()

    val_loss = running_loss / processed_size
    val_score = running_score / processed_size
    return val_loss, val_score
```

In [0]:

```
def fit_epoch(model, train_loader, criterion, metric, op, threshold):
    running_loss = .0
    running_score = .0
    processed_data = len(train_loader)

    model = model.to(device)
    model.train()
    for inputs, real_masks in train_loader:
        inputs = inputs.to(device)
        real_masks = real_masks.to(device)
        op.zero_grad()

        outputs = model(inputs)
        loss = criterion(outputs, real_masks)
        loss.backward()
        op.step()

        pred_masks = torch.sigmoid(outputs) > threshold
        running_score += metric(pred_masks, real_masks).mean().item()
        running_loss += loss.item()

    train_loss = running_loss / processed_data
    train_score = running_score / processed_data
    return train_loss, train_score
```

In [0]:

```
def get_lr(optimizer):
    for param_group in optimizer.param_groups:
        return param_group['lr']
```


In [0]:

```
def train(name_model, model, opt, loss_fn, metric, epochs, data_tr, data_val, threshold=.7):
    X_val, Y_val = next(iter(data_val))
    history = list()
    best_score = 0.
    scheduler = torch.optim.lr_scheduler.StepLR(opt, step_size=35, gamma=0.1)

    for epoch in range(epochs):
        tic = time()
        print('* Epoch %d/%d' % (epoch+1, epochs))

        train_loss, train_score = fit_epoch(model,
                                            data_tr,
                                            loss_fn,
                                            metric,
                                            opt,
                                            threshold
                                            )

        toc = time()
        print('* train_loss: {np.round(train_loss, 4)}, train_score: {np.round(train_score, 4)}')

        val_loss, val_score = eval_epoch(model,
                                         data_val,
                                         loss_fn,
                                         metric,
                                         threshold
                                         )

        print('* val_loss: {np.round(val_loss, 4)}, val_score: {np.round(val_score, 4)}')

        scheduler.step()
        print("* lr:", get_lr(opt))
        history.append((train_loss, train_score, val_loss, val_score))

        if best_score < val_score:
            save_weights(model, f"model_{name_model}.pth")
            best_score = val_score

    # show intermediate results
    model.eval() # testing mode
    Y_hat = model(X_val.to(device)).detach().cpu() # detach and put into cpu

    # Visualize tools
    clear_output(wait=True)
    for k in range(6):
        plt.subplot(4, 6, k+1)
        plt.imshow(np.rollaxis(X_val[k].numpy(), 0, 3), cmap='gray')
        plt.title('Real')
        plt.axis('off')

        plt.subplot(4, 6, k+7)
        plt.imshow(Y_val[k].numpy().squeeze(), cmap='gray')
        plt.title('Real_mask')
        plt.axis('off')

        plt.subplot(4, 6, k+13)
        plt.imshow(Y_hat[k, 0], cmap='gray')
        plt.title('Output')
        plt.axis('off')

        plt.subplot(4, 6, k+19)
        plt.imshow(torch.sigmoid(Y_hat[k, 0]) > threshold, cmap='gray')
        plt.title('Output_mask')
        plt.axis('off')

        info = f'{epoch+1} / {epochs} - train_loss: {np.round(train_loss, 4)}, train_score: {np.round(train_score, 4)}, val_loss: {np.round(val_loss, 4)}, val_score: {np.round(val_score, 4)}, max_val_score: {np.round(best_score, 4)}'
        if (epoch+1) % 10 == 0 and not (epoch+1) == epochs:
            plot_stat(history)
            plt.suptitle(info)
            plt.show()

    load_weights(model, paths=[f"/content/drive/My Drive/Colab Notebooks/Segmentation/weights/model_{name_model}.pth"])
    return history
```

In [0]:

```
def plot_stat(history, figsize=(15, 9)):
    train_loss, train_score, val_loss, val_score = zip(*history)
    max_score_train = np.max(train_score)
    max_score_loss_train = train_loss[np.argmax(train_score)]
    max_score_val = np.max(val_score)
    max_score_loss_val = val_loss[np.argmax(val_score)]
    print(f'max score on train: {max_score_train} with loss: {max_score_loss_train} \nmax score on val: {max_score_val} with loss: {max_score_loss_val} ')
    plt.figure(figsize=figsize)
    plt.subplot(2, 1, 1)
    plt.plot(train_score, label="train_score")
    plt.plot(val_score, label="val_score")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("score")
    plt.grid()
    plt.subplot(2, 1, 2)
    plt.plot(train_loss, label="train_loss")
    plt.plot(val_loss, label="val_loss")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("loss")
    plt.grid()
    plt.show()
```

Инференс [1 балл]

После обучения модели эту функцию можно использовать для прогнозирования сегментации на новых данных:

In [0]:

```
def eval(model, loader, criterion, metric, threshold=.7):
    running_loss = .0
    running_score = .0
    processed_size = len(loader)

    model = model.to(device)
    model.eval()
    for inputs, real_masks in loader:
        inputs = inputs.to(device)
        real_masks = real_masks.to(device)

        with torch.no_grad():
            outputs = model(inputs)
            loss = criterion(outputs, real_masks)
            pred_masks = torch.sigmoid(outputs) > threshold

        running_score += metric(pred_masks, real_masks).mean().item()
        running_loss += loss.item()

    data_loss = running_loss / processed_size
    data_score = running_score / processed_size
    return data_loss, data_score
```

In [0]:

```
def score_model(model, metric, data, threshold=.7):
    model.eval() # testing mode
    scores = 0
    for X_batch, Y_label in data:
        Y_pred = torch.sigmoid(model(X_batch.to(device)))
        scores += metric(Y_pred > threshold, Y_label.to(device)).mean().item()
    return scores/len(data)
```

Основной момент: обучение

In [0]:


```
lr = 1e-4
max_epochs = 40
```

In [28]:

```
seed_init()
clear_memory()
model_bce = SegNet().to(device)

opt = optim.Adam(model_bce.parameters(), lr=lr)
history_bce = train('bce', model_bce, opt, bce_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.7153, train_score: 0.667, val_loss: 0.718, val_score: 0.654, max_val_score: 0.668



In [29]:

```
score_model(model_bce, iou_pytorch, data_val)
```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

"See the documentation of nn.Upsample for details.".format(mode))

Out[29]:


0.6679999828338623

In [30]:

```
plot_stat(history_bce)
```

max score on train: 0.6670000404119492 with loss: 0.7152881622314453

max score on val: 0.6680000126361847 with loss: 0.7214172184467316



Дополнительные функции лосса [2 балла]

В данном разделе вам потребуется имплементировать две функции потерь: DICE и Focal loss.

1. Dice coefficient: Учитывая две маски X и Y , общая метрика для измерения расстояния между этими двумя масками задается следующим образом:

$$D(X,Y)=\frac{2|X\cap Y|}{|X|+|Y|}$$

Эта функция не является дифференцируемой, но это необходимое свойство для градиентного спуска. В данном случае мы можем приблизить его с помощью:

$$\mathcal{L}_D(X,Y) = 1 - \frac{1}{256 \times 256} \times \sum_i \frac{2X_iY_i}{X_i+Y_i}$$

// Не забудьте подумать о численной неустойчивости.

In [0]:

```
def dice_loss(y_pred, y_real, eps=1e-8):
    y_pred, y_real = torch.sigmoid(y_pred.view(-1)), y_real.view(-1)
    num = 2. * y_pred * y_real
    den = y_pred + y_real + eps
    res = torch.mean(num / den)
    return 1. - res
```

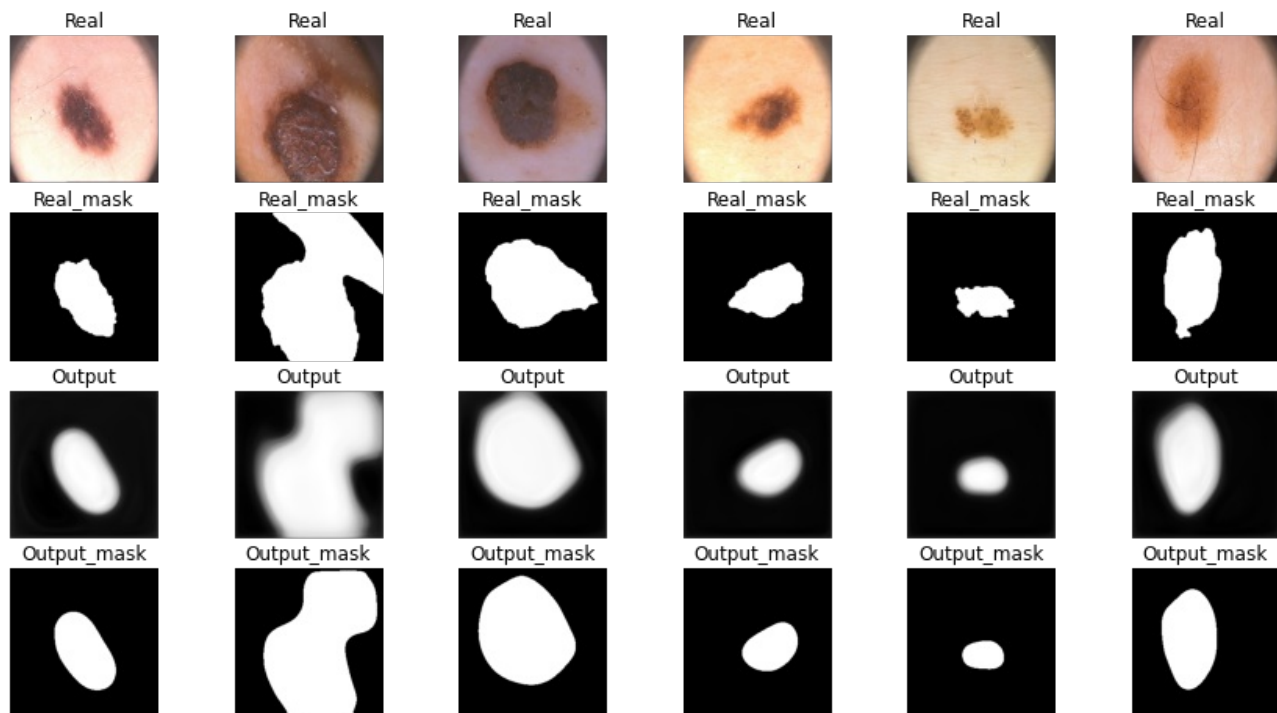
Проводим тестирование:

In [32]:

```
seed_init()
clear_memory()
model_dice = SegNet().to(device)

optimaiz = optim.Adam(model_dice.parameters(), lr=lr)
history_dice = train('dice', model_dice, optimaiz, dice_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.7104, train_score: 0.61, val_loss: 0.7351, val_score: 0.612, max_val_score: 0.642



In [33]:

```
score_model(model_dice, iou_pytorch, data_val)
```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

"See the documentation of nn.Upsample for details.".format(mode))

Out[33]:


0.6420000195503235

In [34]:

```
plot_stat(history_dice)
```

max score on train: 0.6449999958276749 with loss: 0.7093066871166229

max score on val: 0.6420000195503235 with loss: 0.7307058572769165



2. Focal loss: (<https://arxiv.org/pdf/1708.02002.pdf>)

Окей, мы уже с вами умеем делать BCE loss:

$$\mathcal{L}_{\text{BCE}}(y, \hat{y}) = -\sum_i \left[y_i \log(\sigma(\hat{y}_i)) + (1-y_i) \log(1-\sigma(\hat{y}_i)) \right]$$

Проблема с этой потерей заключается в том, что она имеет тенденцию приносить пользу классу **большинства** (фоновому) по отношению к классу **меньшинства** (переднему). Поэтому обычно применяются весовые коэффициенты к каждому классу:

$$\mathcal{L}_{\text{wBCE}}(y, \hat{y}) = -\sum_i \alpha_i \left[y_i \log(\sigma(\hat{y}_i)) + (1-y_i) \log(1-\sigma(\hat{y}_i)) \right]$$

Традиционно вес α_i определяется как обратная частота класса этого пикселя i , так что наблюдения миноритарного класса весят больше по отношению к классу большинства.

Еще одним недавним дополнением является взвешенный пиксельный вариант, которая взвешивает каждый пиксель по степени уверенности, которую мы имеем в предсказании этого пикселя.

$$\mathcal{L}_{\text{focal}}(y, \hat{y}) = -\sum_i \left[\left(1 - \sigma(\hat{y}_i) \right)^\gamma y_i \log(\sigma(\hat{y}_i)) + (1-y_i) \log(1-\sigma(\hat{y}_i)) \right]$$

Зафиксируем значение $\gamma=2$.

In [0]:

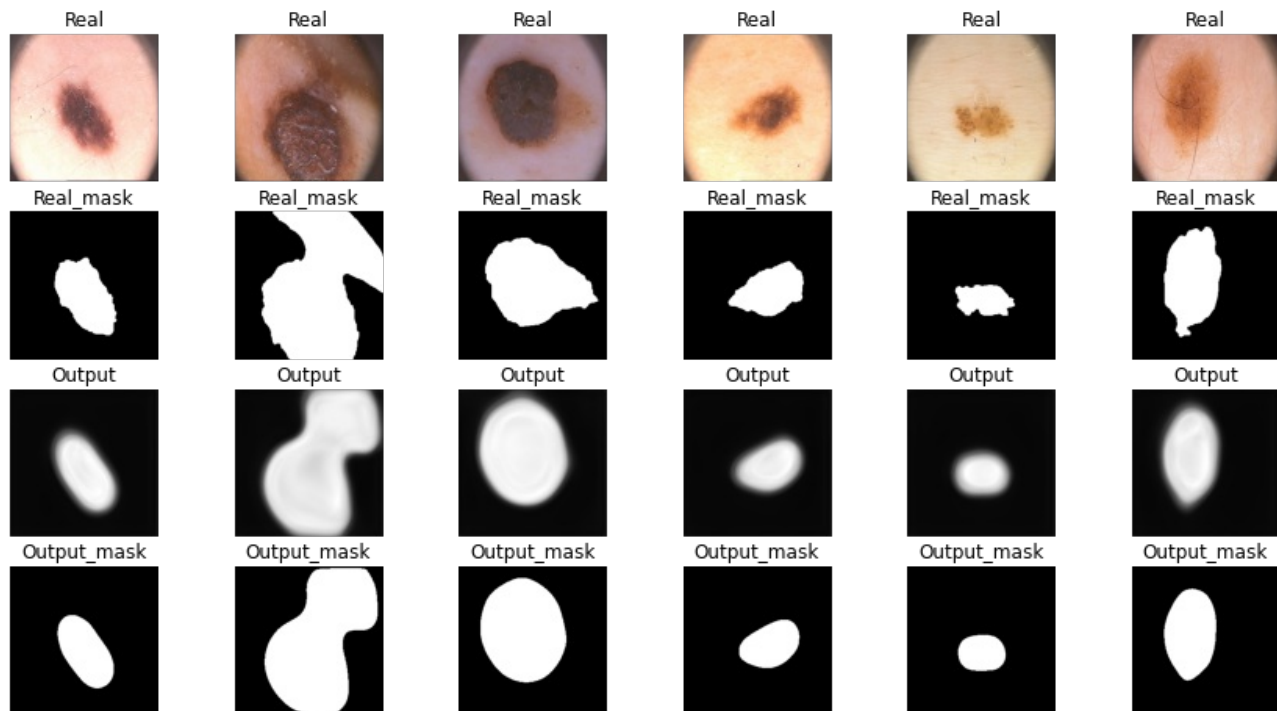
```
def focal_loss(y_pred, y_real, alpha=1, gamma=2, eps=1e-8, bce_loss=bce_loss):
    y_pred, y_real = y_pred.view(y_pred.shape[0], -1), y_real.view(y_pred.shape[0], -1)
    bce = bce_loss(y_pred, y_real, reduce=False) # F.binary_cross_entropy_with_logits(y_pred, y_real, reduce=False)
    pt = torch.exp(-bce)
    loss = alpha * ((1 - pt) ** gamma) * bce
    # y_pred = hint: torch.clamp
    return loss.mean()
```

In [36]:

```
seed_init()
clear_memory()
model_focal = SegNet().to(device)

optimaiz = optim.Adam(model_focal.parameters(), lr=lr)
history_focal = train('focal', model_focal, optimaiz, focal_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.2157, train_score: 0.643, val_loss: 0.2164, val_score: 0.602, max_val_score: 0.63



In [37]:

```
score_model(model_focal, iou_pytorch, data_val)
```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

"See the documentation of nn.Upsample for details.".format(mode))

Out[37]:


0.6299999952316284

In [38]:

```
plot_stat(history_focal)
```

max score on train: 0.6470000147819519 with loss: 0.21528907120227814

max score on val: 0.6300000250339508 with loss: 0.21781566739082336



[BONUS] Мир сегментационных лоссов [5 баллов]

В данном блоке предлагаю написать вам 1 функцию потерь самостоятельно. Для этого необходимо прочитать статью и имплементировать ее. Кроме того провести численное сравнение с предыдущими функциями. Какие варианты?

1) Можно учесть Total Variation 2) Lova 3) BCE но с Soft Targets (что-то типа label-smoothing для многослассовой классификации) 4) Tversky loss 5) Любой другой

- [Physiological Inspired Deep Neural Networks for Emotion Recognition \(https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=8472816&tag=1\)](https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=8472816&tag=1). IEEE Access, 6, 53930-53943.
- [Boundary loss for highly unbalanced segmentation \(https://arxiv.org/abs/1812.07032\)](https://arxiv.org/abs/1812.07032)
- [Tversky loss function for image segmentation using 3D fully convolutional deep networks \(https://arxiv.org/abs/1706.05721\)](https://arxiv.org/abs/1706.05721)
- [Correlation Maximized Structural Similarity Loss for Semantic Segmentation \(https://arxiv.org/abs/1910.08711\)](https://arxiv.org/abs/1910.08711)
- [Topology-Preserving Deep Image Segmentation \(https://papers.nips.cc/paper/8803-topology-preserving-deep-image-segmentation\)](https://papers.nips.cc/paper/8803-topology-preserving-deep-image-segmentation)

Exercise: Add the total variation term to the loss.

In [0]:


```
#<TODO>
def tversky_loss(beta, eps = 1e-8):
    def loss(y_pred, y_real):
        y_pred, y_real = torch.sigmoid(y_pred.view(-1)), y_real.view(-1)
        numerator = torch.sum(y_real * y_pred)
        denominator = y_real * y_pred + beta * (1. - y_real) * y_pred + (1. - beta) * y_real * (1. - y_pred)
        return 1. - numerator / (torch.sum(denominator) + eps)
    return loss
```


In [40]:

```
seed_init()
clear_memory()
model_tversky = SegNet().to(device)
tversky_loss_07 = tversky_loss(beta=.3)

optima = optim.Adam(model_tversky.parameters(), lr=lr)
history_tversky = train('tversky', model_tversky, optima, tversky_loss_07, iou_pytorch, max_epochs, data_tr,
data_val)
```

40 / 40 - train_loss: 0.3305, train_score: 0.637, val_loss: 0.3476, val_score: 0.652, max_val_score: 0.658



In [41]:

```
score_model(model_tversky, iou_pytorch, data_val)
```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.

"See the documentation of nn.Upsample for details.".format(mode))

Out[41]:


0.6579999923706055

In [42]:

```
plot_stat(history_tversky)
```

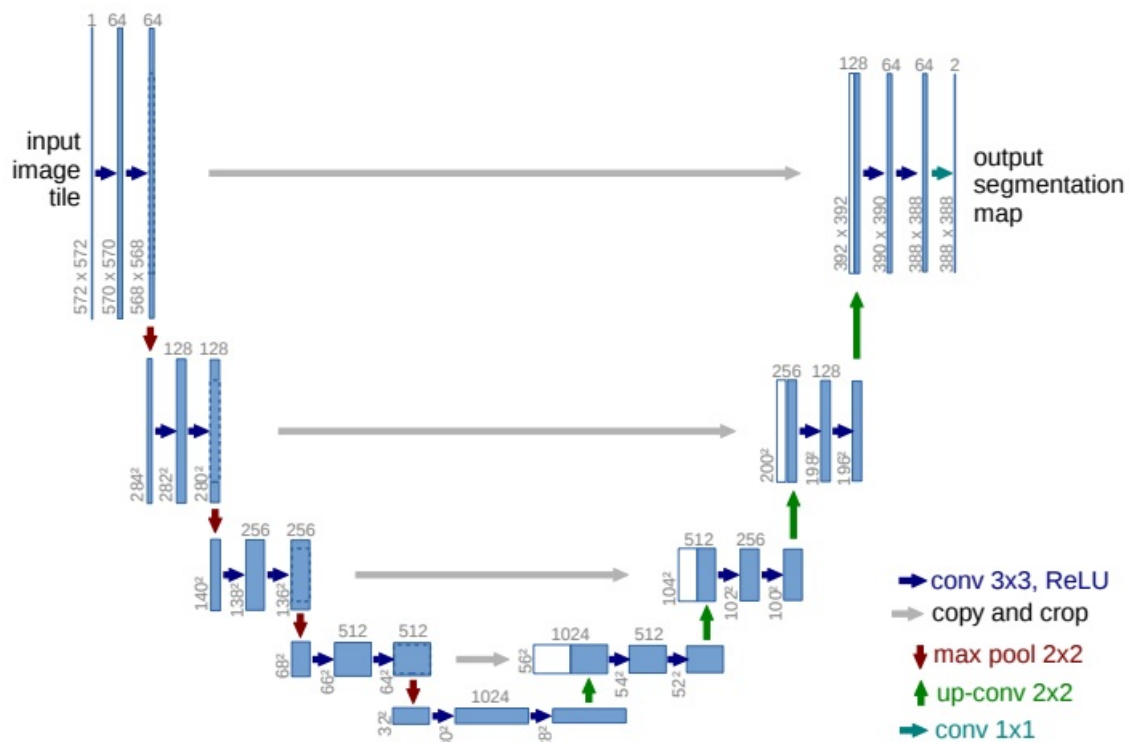
max score on train: 0.6500000059604645 with loss: 0.3270793706178665

max score on val: 0.6579999923706055 with loss: 0.3447708189487457



U-Net [2 балла]

U-Net (<https://arxiv.org/abs/1505.04597>) это архитектура нейронной сети, которая получает изображение и выводит его. Первоначально он был задуман для семантической сегментации (как мы ее будем использовать), но он настолько успешен, что с тех пор используется в других контекстах. Учитывая медицинское изображение, он выводит изображение в оттенках серого, представляющее вероятность того, что каждый пиксель является интересующей областью.



У нас в архитектуре все так же существует энкодер и декодер, как в **SegNet**, но отличительной особенностью данной модели являются skip-connections. Элементы соединяющие части декодера и энкодера. То есть для того чтобы передать на вход декодера тензор, мы конкатенируем симметричный выход с энкодера и выход предыдущего слоя декодера.

- Ronneberger, Olaf, Philipp Fischer, and Thomas Brox. "U-Net: Convolutional networks for biomedical image segmentation." <https://arxiv.org/pdf/1505.04597.pdf> International Conference on Medical image computing and computer-assisted intervention. Springer, Cham, 2015.

In [0]:

```
class UNet(nn.Module):
    def __init__(self):
        super().__init__()

        # encoder (downsampling)
        self.enc_conv0 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
        )

        self.pool0 = nn.MaxPool2d(kernel_size=2, stride=2) # 256 -> 128
        self.enc_conv1 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
        )

        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2) # 128 -> 64
        self.enc_conv2 = nn.Sequential(
            nn.Conv2d(64, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.3),
            nn.Conv2d(128, 128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
        )

        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2) # 64 -> 32
        self.enc_conv3 = nn.Sequential(
            nn.Conv2d(128, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.3),
            nn.Conv2d(256, 256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
        )

        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2) # 32 -> 16

        # bottleneck
        self.bottleneck_conv = nn.Sequential(
            nn.Conv2d(256, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
            nn.Dropout2d(p=0.3),
            nn.Conv2d(512, 512, kernel_size=3, padding=1),
            nn.BatchNorm2d(512),
            nn.ReLU(inplace=True),
        )

        # decoder (upsampling)
        self.upsample0 = nn.Upsample(scale_factor=2, mode="bilinear") # 16 -> 32
        self.dec_conv0 = nn.Sequential(
```

```

nn.Conv2d(512, 256, kernel_size=3, padding=1),

nn.BatchNorm2d(256),
nn.ReLU(inplace=True),
nn.Dropout2d(p=0.3),
nn.Conv2d(256, 256, kernel_size=3, padding=1),
nn.BatchNorm2d(256),
nn.ReLU(inplace=True),

)
# *2 - because concatenation
self.upsample1 = nn.Upsample(scale_factor=2, mode="bilinear") # 32 -> 64
self.dec_conv1 = nn.Sequential(
    nn.Conv2d(256*2, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(128, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),

)
# *2 - because concatenation
self.upsample2 = nn.Upsample(scale_factor=2, mode="bilinear") # 64 -> 128
self.dec_conv2 = nn.Sequential(
    nn.Conv2d(128*2, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),
    nn.Conv2d(64, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),

)
# *2 - because concatenation
self.upsample3 = nn.Upsample(scale_factor=2, mode="bilinear") # 128 -> 256
self.dec_conv3 = nn.Sequential(
    nn.Conv2d(64*2, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.ReLU(inplace=True),
    nn.Conv2d(32, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),
    nn.ReLU(inplace=True),

)
# *2 - because concatenation
self.dec_final = nn.Sequential(nn.Conv2d(32*2, 1, kernel_size=3, padding=1),
                                nn.BatchNorm2d(1)
                                )

def forward(self, x):
    # encoder
    e0 = self.pool0(self.enc_conv0(x))
    e1 = self.pool1(self.enc_conv1(e0))
    e2 = self.pool2(self.enc_conv2(e1))
    e3 = self.pool3(self.enc_conv3(e2))

    # bottleneck
    b = self.bottleneck_conv(e3)

    # decoder
    # print('enc_part: ', self.enc_conv3(e2).shape, '+', 'dec_part: ', self.dec_conv0(self.upsample0(b)).
shape)
    d0 = torch.cat((self.enc_conv3(e2), self.dec_conv0(self.upsample0(b))), 1)
    # print('enc_part: ', self.enc_conv2(e1).shape, '+', 'dec_part: ', self.dec_conv1(self.upsample1(d0))
.shape)
    d1 = torch.cat((self.enc_conv2(e1), self.dec_conv1(self.upsample1(d0))), 1)
    # print('enc_part: ', self.enc_conv1(e0).shape, '+', 'dec_part: ', self.dec_conv2(self.upsample2(d1))
.shape)
    d2 = torch.cat((self.enc_conv1(e0), self.dec_conv2(self.upsample2(d1))), 1)
    # print('enc_part: ', self.enc_conv0(x).shape, '+', 'dec_part: ', self.dec_conv3(self.upsample3(d2)).
shape)
    d3 = torch.cat((self.enc_conv0(x), self.dec_conv3(self.upsample3(d2))), 1)
    # print(d3.shape)
    dfin = self.dec_final(d3) # no activation
    # print(dfin.shape)
    return dfin

```

In [44]:

```
summary(UNet().to(device), input_size=(3, 256, 256))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	9,248
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
MaxPool2d-7	[-1, 32, 128, 128]	0
Conv2d-8	[-1, 64, 128, 128]	18,496
BatchNorm2d-9	[-1, 64, 128, 128]	128
ReLU-10	[-1, 64, 128, 128]	0
Conv2d-11	[-1, 64, 128, 128]	36,928
BatchNorm2d-12	[-1, 64, 128, 128]	128
ReLU-13	[-1, 64, 128, 128]	0
MaxPool2d-14	[-1, 64, 64, 64]	0
Conv2d-15	[-1, 128, 64, 64]	73,856
BatchNorm2d-16	[-1, 128, 64, 64]	256
ReLU-17	[-1, 128, 64, 64]	0
Dropout2d-18	[-1, 128, 64, 64]	0
Conv2d-19	[-1, 128, 64, 64]	147,584
BatchNorm2d-20	[-1, 128, 64, 64]	256
ReLU-21	[-1, 128, 64, 64]	0
MaxPool2d-22	[-1, 128, 32, 32]	0
Conv2d-23	[-1, 256, 32, 32]	295,168
BatchNorm2d-24	[-1, 256, 32, 32]	512
ReLU-25	[-1, 256, 32, 32]	0
Dropout2d-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 256, 32, 32]	590,080
BatchNorm2d-28	[-1, 256, 32, 32]	512
ReLU-29	[-1, 256, 32, 32]	0
MaxPool2d-30	[-1, 256, 16, 16]	0
Conv2d-31	[-1, 512, 16, 16]	1,180,160
BatchNorm2d-32	[-1, 512, 16, 16]	1,024
ReLU-33	[-1, 512, 16, 16]	0
Dropout2d-34	[-1, 512, 16, 16]	0
Conv2d-35	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-36	[-1, 512, 16, 16]	1,024
ReLU-37	[-1, 512, 16, 16]	0
Conv2d-38	[-1, 256, 32, 32]	295,168
BatchNorm2d-39	[-1, 256, 32, 32]	512
ReLU-40	[-1, 256, 32, 32]	0
Dropout2d-41	[-1, 256, 32, 32]	0
Conv2d-42	[-1, 256, 32, 32]	590,080
BatchNorm2d-43	[-1, 256, 32, 32]	512
ReLU-44	[-1, 256, 32, 32]	0
Upsample-45	[-1, 512, 32, 32]	0
Conv2d-46	[-1, 256, 32, 32]	1,179,904
BatchNorm2d-47	[-1, 256, 32, 32]	512
ReLU-48	[-1, 256, 32, 32]	0
Dropout2d-49	[-1, 256, 32, 32]	0
Conv2d-50	[-1, 256, 32, 32]	590,080
BatchNorm2d-51	[-1, 256, 32, 32]	512
ReLU-52	[-1, 256, 32, 32]	0
Conv2d-53	[-1, 128, 64, 64]	73,856
BatchNorm2d-54	[-1, 128, 64, 64]	256
ReLU-55	[-1, 128, 64, 64]	0
Dropout2d-56	[-1, 128, 64, 64]	0
Conv2d-57	[-1, 128, 64, 64]	147,584
BatchNorm2d-58	[-1, 128, 64, 64]	256
ReLU-59	[-1, 128, 64, 64]	0
Upsample-60	[-1, 512, 64, 64]	0
Conv2d-61	[-1, 128, 64, 64]	589,952
BatchNorm2d-62	[-1, 128, 64, 64]	256
ReLU-63	[-1, 128, 64, 64]	0
Dropout2d-64	[-1, 128, 64, 64]	0
Conv2d-65	[-1, 128, 64, 64]	147,584
BatchNorm2d-66	[-1, 128, 64, 64]	256
ReLU-67	[-1, 128, 64, 64]	0
Conv2d-68	[-1, 64, 128, 128]	18,496
BatchNorm2d-69	[-1, 64, 128, 128]	128
ReLU-70	[-1, 64, 128, 128]	0
Conv2d-71	[-1, 64, 128, 128]	36,928
BatchNorm2d-72	[-1, 64, 128, 128]	128
ReLU-73	[-1, 64, 128, 128]	0
Upsample-74	[-1, 256, 128, 128]	0
Conv2d-75	[-1, 64, 128, 128]	147,520
BatchNorm2d-76	[-1, 64, 128, 128]	128
ReLU-77	[-1, 64, 128, 128]	0

Conv2d-78	[-1, 64, 128, 128]	36,928
BatchNorm2d-79	[-1, 64, 128, 128]	128
ReLU-80	[-1, 64, 128, 128]	0
Conv2d-81	[-1, 32, 256, 256]	896
BatchNorm2d-82	[-1, 32, 256, 256]	64
ReLU-83	[-1, 32, 256, 256]	0
Conv2d-84	[-1, 32, 256, 256]	9,248
BatchNorm2d-85	[-1, 32, 256, 256]	64
ReLU-86	[-1, 32, 256, 256]	0
Upsample-87	[-1, 128, 256, 256]	0
Conv2d-88	[-1, 32, 256, 256]	36,896
BatchNorm2d-89	[-1, 32, 256, 256]	64
ReLU-90	[-1, 32, 256, 256]	0
Conv2d-91	[-1, 32, 256, 256]	9,248
BatchNorm2d-92	[-1, 32, 256, 256]	64
ReLU-93	[-1, 32, 256, 256]	0
Conv2d-94	[-1, 1, 256, 256]	577
BatchNorm2d-95	[-1, 1, 256, 256]	2

```

=====
Total params: 8,630,979
Trainable params: 8,630,979
Non-trainable params: 0
-----

```

```

Input size (MB): 0.75
Forward/backward pass size (MB): 689.50
Params size (MB): 32.92
Estimated Total Size (MB): 723.17
-----

```

```

/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsamp
ling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify
align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for det
ails.

```

```

"See the documentation of nn.Upsample for details.".format(mode))

```

In [0]:

```

seed_init()
clear_memory()

UNET_MODEL = UNet().to(device)
op = optim.Adam(UNET_MODEL.parameters(), lr=lr)

```

In [46]:

```


history_UNET = train("UNET", UNET_MODEL, op, bce_loss, iou_pytorch, max_epochs, data_tr, data_val)

```

```

40 / 40 - train_loss: 0.7128, train_score: 0.702, val_loss: 0.7179, val_score: 0.694, max_val_score: 0.704

```



```
plot_stat(history_unet)
```

The figure consists of two line plots sharing a common x-axis representing 'epochs' from 0 to 40.

The top plot shows 'score' on the y-axis (0.0 to 0.7). The 'train_score' (blue line) starts at 0.2 and increases steadily to approximately 0.7. The 'val_score' (orange line) starts at 0.0, remains low until epoch 4, then increases sharply to meet the training score around epoch 15, and continues to rise slightly with some fluctuations, reaching approximately 0.7 by epoch 40.

The bottom plot shows 'loss' on the y-axis (0.72 to 0.80). The 'train_loss' (blue line) starts at approximately 0.755 and decreases to about 0.715. The 'val_loss' (orange line) starts at approximately 0.85 and decreases sharply to about 0.725 by epoch 10, then continues to decrease slowly, stabilizing around 0.718 by epoch 40.

```
score_model(unet_model, iou_pytorch, data_val)
```

```
"See the documentation of nn.Upsample for details.".format(mode))
```

0.7039999663829803

Max-Pooling for the downsampling and **nearest-neighbor Upsampling** for the upsampling.

```
conv = nn.Conv2d(3, 64, 3, padding=1)
pool = nn.MaxPool2d(3, 2, padding=1)
```

```
upsample = nn.Upsample(32)
conv = nn.Conv2d(64, 64, 3, padding=1)
```

```
class UNet2(nn.Module):
    def __init__(self, near_neigh=True):
        super().__init__()

        # encoder (downsampling)
        self.enc_conv0 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, kernel_size=3, padding=1)
```

```

        nn.Conv2d(32, 32, kernel_size=3, padding=1),
        nn.BatchNorm2d(32),
        nn.ReLU(inplace=True),

    )

    self.pool0 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1) if near_neigh else nn.Conv2d(32, 32, k
kernel_size=2, stride=2) # 256 -> 128
    self.enc_conv1 = nn.Sequential(
        nn.Conv2d(32, 64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(inplace=True),
        nn.Conv2d(64, 64, kernel_size=3, padding=1),
        nn.BatchNorm2d(64),
        nn.ReLU(inplace=True),

    )

    self.pool1 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1) if near_neigh else nn.Conv2d(64, 64, k
kernel_size=2, stride=2) # 128 -> 64
    self.enc_conv2 = nn.Sequential(
        nn.Conv2d(64, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),
        nn.Dropout2d(p=0.3),
        nn.Conv2d(128, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),

    )

    self.pool2 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1) if near_neigh else nn.Conv2d(128, 128,
kernel_size=2, stride=2) # 64 -> 32
    self.enc_conv3 = nn.Sequential(
        nn.Conv2d(128, 256, kernel_size=3, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),
        nn.Dropout2d(p=0.3),
        nn.Conv2d(256, 256, kernel_size=3, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),

    )

    self.pool3 = nn.MaxPool2d(kernel_size=3, stride=2, padding=1) if near_neigh else nn.Conv2d(256, 256,
kernel_size=2, stride=2) # 32 -> 16

    # bottleneck
    self.bottleneck_conv = nn.Sequential(
        nn.Conv2d(256, 512, kernel_size=3, padding=1),
        nn.BatchNorm2d(512),
        nn.ReLU(inplace=True),
        nn.Dropout2d(p=0.3),
        nn.Conv2d(512, 512, kernel_size=3, padding=1),
        nn.BatchNorm2d(512),
        nn.ReLU(inplace=True),

    )

    # decoder (upsampling)
    self.upsample0 = nn.Upsample(scale_factor=2, mode='nearest') if near_neigh else nn.ConvTranspose2d(
512, 512, kernel_size=2, stride=2) # 16 -> 32 # torch.nn.ConvTranspose2d(in_channels=512, out_channels=256,
kernel_size=3, stride=2, padding=1, output_padding=1)
    self.dec_conv0 = nn.Sequential(
        nn.Conv2d(512, 256, kernel_size=3, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),
        nn.Dropout2d(p=0.3),
        nn.Conv2d(256, 256, kernel_size=3, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),

    )

    # *2 - because concatenation
    self.upsample1 = nn.Upsample(scale_factor=2, mode='nearest') if near_neigh else nn.ConvTranspose2d(
256*2, 256*2, kernel_size=2, stride=2) # 32 -> 64
    self.dec_conv1 = nn.Sequential(
        nn.Conv2d(256*2, 128, kernel_size=3, padding=1),
        nn.BatchNorm2d(128),
        nn.ReLU(inplace=True),

```



```

nn.Dropout2d(p=0.3),
nn.Conv2d(128, 128, kernel_size=3, padding=1),
nn.BatchNorm2d(128),
nn.ReLU(inplace=True),
)
# *2 - because concatenation
self.upsample2 = nn.Upsample(scale_factor=2, mode='nearest') if near_neigh else nn.ConvTranspose2d(
128*2, 128*2, kernel_size=2, stride=2) # 64 -> 128
self.dec_conv2 = nn.Sequential(
nn.Conv2d(128*2, 64, kernel_size=3, padding=1),
nn.BatchNorm2d(64),
nn.ReLU(inplace=True),
nn.Conv2d(64, 64, kernel_size=3, padding=1),
nn.BatchNorm2d(64),
nn.ReLU(inplace=True),
)
# *2 - because concatenation
self.upsample3 = nn.Upsample(scale_factor=2, mode='nearest') if near_neigh else nn.ConvTranspose2d(
64*2, 64*2, kernel_size=2, stride=2) # 128 -> 256
self.dec_conv3 = nn.Sequential(
nn.Conv2d(64*2, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(inplace=True),
nn.Conv2d(32, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(inplace=True),
)
# *2 - because concatenation
self.dec_final = nn.Sequential(nn.Conv2d(32*2, 1, kernel_size=3, padding=1),
nn.BatchNorm2d(1)
)

def forward(self, x):
# encoder
e0 = self.pool0(self.enc_conv0(x))
e1 = self.pool1(self.enc_conv1(e0))
e2 = self.pool2(self.enc_conv2(e1))
e3 = self.pool3(self.enc_conv3(e2))

# bottleneck
b = self.bottleneck_conv(e3)

# decoder
# print('enc_part: ', self.enc_conv3(e2).shape, '+', 'dec_part: ', self.dec_conv0(self.upsample0(b)).
shape)
d0 = torch.cat((self.enc_conv3(e2), self.dec_conv0(self.upsample0(b))), 1)
# print('enc_part: ', self.enc_conv2(e1).shape, '+', 'dec_part: ', self.dec_conv1(self.upsample1(d0))
.shape)
d1 = torch.cat((self.enc_conv2(e1), self.dec_conv1(self.upsample1(d0))), 1)
# print('enc_part: ', self.enc_conv1(e0).shape, '+', 'dec_part: ', self.dec_conv2(self.upsample2(d1))
.shape)
d2 = torch.cat((self.enc_conv1(e0), self.dec_conv2(self.upsample2(d1))), 1)
# print('enc_part: ', self.enc_conv0(x).shape, '+', 'dec_part: ', self.dec_conv3(self.upsample3(d2)).
shape)
d3 = torch.cat((self.enc_conv0(x), self.dec_conv3(self.upsample3(d2))), 1)
# print(d3.shape)
dfin = self.dec_final(d3) # no activation
# print(dfin.shape)
return dfin

```

In [50]:

```
summary(UNet2().to(device), input_size=(3, 256, 256))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	9,248
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
MaxPool2d-7	[-1, 32, 128, 128]	0

Conv2d-8	[-1, 64, 128, 128]	18,496
BatchNorm2d-9	[-1, 64, 128, 128]	128
ReLU-10	[-1, 64, 128, 128]	0
Conv2d-11	[-1, 64, 128, 128]	36,928
BatchNorm2d-12	[-1, 64, 128, 128]	128
ReLU-13	[-1, 64, 128, 128]	0
MaxPool2d-14	[-1, 64, 64, 64]	0
Conv2d-15	[-1, 128, 64, 64]	73,856
BatchNorm2d-16	[-1, 128, 64, 64]	256
ReLU-17	[-1, 128, 64, 64]	0
Dropout2d-18	[-1, 128, 64, 64]	0
Conv2d-19	[-1, 128, 64, 64]	147,584
BatchNorm2d-20	[-1, 128, 64, 64]	256
ReLU-21	[-1, 128, 64, 64]	0
MaxPool2d-22	[-1, 128, 32, 32]	0
Conv2d-23	[-1, 256, 32, 32]	295,168
BatchNorm2d-24	[-1, 256, 32, 32]	512
ReLU-25	[-1, 256, 32, 32]	0
Dropout2d-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 256, 32, 32]	590,080
BatchNorm2d-28	[-1, 256, 32, 32]	512
ReLU-29	[-1, 256, 32, 32]	0
MaxPool2d-30	[-1, 256, 16, 16]	0
Conv2d-31	[-1, 512, 16, 16]	1,180,160
BatchNorm2d-32	[-1, 512, 16, 16]	1,024
ReLU-33	[-1, 512, 16, 16]	0
Dropout2d-34	[-1, 512, 16, 16]	0
Conv2d-35	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-36	[-1, 512, 16, 16]	1,024
ReLU-37	[-1, 512, 16, 16]	0
Conv2d-38	[-1, 256, 32, 32]	295,168
BatchNorm2d-39	[-1, 256, 32, 32]	512
ReLU-40	[-1, 256, 32, 32]	0
Dropout2d-41	[-1, 256, 32, 32]	0
Conv2d-42	[-1, 256, 32, 32]	590,080
BatchNorm2d-43	[-1, 256, 32, 32]	512
ReLU-44	[-1, 256, 32, 32]	0
Upsample-45	[-1, 512, 32, 32]	0
Conv2d-46	[-1, 256, 32, 32]	1,179,904
BatchNorm2d-47	[-1, 256, 32, 32]	512
ReLU-48	[-1, 256, 32, 32]	0
Dropout2d-49	[-1, 256, 32, 32]	0
Conv2d-50	[-1, 256, 32, 32]	590,080
BatchNorm2d-51	[-1, 256, 32, 32]	512
ReLU-52	[-1, 256, 32, 32]	0
Conv2d-53	[-1, 128, 64, 64]	73,856
BatchNorm2d-54	[-1, 128, 64, 64]	256
ReLU-55	[-1, 128, 64, 64]	0
Dropout2d-56	[-1, 128, 64, 64]	0
Conv2d-57	[-1, 128, 64, 64]	147,584
BatchNorm2d-58	[-1, 128, 64, 64]	256
ReLU-59	[-1, 128, 64, 64]	0
Upsample-60	[-1, 512, 64, 64]	0
Conv2d-61	[-1, 128, 64, 64]	589,952
BatchNorm2d-62	[-1, 128, 64, 64]	256
ReLU-63	[-1, 128, 64, 64]	0
Dropout2d-64	[-1, 128, 64, 64]	0
Conv2d-65	[-1, 128, 64, 64]	147,584
BatchNorm2d-66	[-1, 128, 64, 64]	256
ReLU-67	[-1, 128, 64, 64]	0
Conv2d-68	[-1, 64, 128, 128]	18,496
BatchNorm2d-69	[-1, 64, 128, 128]	128
ReLU-70	[-1, 64, 128, 128]	0
Conv2d-71	[-1, 64, 128, 128]	36,928
BatchNorm2d-72	[-1, 64, 128, 128]	128
ReLU-73	[-1, 64, 128, 128]	0
Upsample-74	[-1, 256, 128, 128]	0
Conv2d-75	[-1, 64, 128, 128]	147,520
BatchNorm2d-76	[-1, 64, 128, 128]	128
ReLU-77	[-1, 64, 128, 128]	0
Conv2d-78	[-1, 64, 128, 128]	36,928
BatchNorm2d-79	[-1, 64, 128, 128]	128
ReLU-80	[-1, 64, 128, 128]	0
Conv2d-81	[-1, 32, 256, 256]	896
BatchNorm2d-82	[-1, 32, 256, 256]	64
ReLU-83	[-1, 32, 256, 256]	0
Conv2d-84	[-1, 32, 256, 256]	9,248
BatchNorm2d-85	[-1, 32, 256, 256]	64
ReLU-86	[-1, 32, 256, 256]	0
Upsample-87	[-1, 128, 256, 256]	0

Conv2d-88	[-1, 32, 256, 256]	36,896
BatchNorm2d-89	[-1, 32, 256, 256]	64
ReLU-90	[-1, 32, 256, 256]	0
Conv2d-91	[-1, 32, 256, 256]	9,248
BatchNorm2d-92	[-1, 32, 256, 256]	64
ReLU-93	[-1, 32, 256, 256]	0
Conv2d-94	[-1, 1, 256, 256]	577
BatchNorm2d-95	[-1, 1, 256, 256]	2

```

=====
Total params: 8,630,979
Trainable params: 8,630,979
Non-trainable params: 0
-----

```

```

-----
Input size (MB): 0.75
Forward/backward pass size (MB): 689.50
Params size (MB): 32.92
Estimated Total Size (MB): 723.17
-----

```

In [51]:


```

seed_init()
clear_memory()

UNET2_1_MODEL = UNet2().to(device)
op = optim.Adam(UNET2_1_MODEL.parameters(), lr=lr)
history_UNET2_1 = train('UNET2_1', UNET2_1_MODEL, op, bce_loss, iou_pytorch, max_epochs, data_tr, data_val)

```

40 / 40 - train_loss: 0.7117, train_score: 0.714, val_loss: 0.7176, val_score: 0.694, max_val_score: 0.718



In [52]:

```
score_model(UNET2_1_MODEL, iou_pytorch, data_val)
```

Out[52]:


0.7179999649524689

In [53]:

```
plot_stat(history_unet_2_1)
```

max score on train: 0.757999986410141 with loss: 0.7112099975347519

max score on val: 0.7179999947547913 with loss: 0.720087319612503



In [54]:

```
summary(UNet2(near_neigh=False).to(device), input_size=(3, 256, 256))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	9,248
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
Conv2d-7	[-1, 32, 128, 128]	4,128
Conv2d-8	[-1, 64, 128, 128]	18,496
BatchNorm2d-9	[-1, 64, 128, 128]	128
ReLU-10	[-1, 64, 128, 128]	0
Conv2d-11	[-1, 64, 128, 128]	36,928
BatchNorm2d-12	[-1, 64, 128, 128]	128
ReLU-13	[-1, 64, 128, 128]	0
Conv2d-14	[-1, 64, 64, 64]	16,448
Conv2d-15	[-1, 128, 64, 64]	73,856
BatchNorm2d-16	[-1, 128, 64, 64]	256
ReLU-17	[-1, 128, 64, 64]	0
Dropout2d-18	[-1, 128, 64, 64]	0
Conv2d-19	[-1, 128, 64, 64]	147,584
BatchNorm2d-20	[-1, 128, 64, 64]	256
ReLU-21	[-1, 128, 64, 64]	0
Conv2d-22	[-1, 128, 32, 32]	65,664
Conv2d-23	[-1, 256, 32, 32]	295,168
BatchNorm2d-24	[-1, 256, 32, 32]	512
ReLU-25	[-1, 256, 32, 32]	0
Dropout2d-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 256, 32, 32]	590,080
BatchNorm2d-28	[-1, 256, 32, 32]	512
ReLU-29	[-1, 256, 32, 32]	0
Conv2d-30	[-1, 256, 16, 16]	262,400
Conv2d-31	[-1, 512, 16, 16]	1,180,160
BatchNorm2d-32	[-1, 512, 16, 16]	1,024
ReLU-33	[-1, 512, 16, 16]	0
Dropout2d-34	[-1, 512, 16, 16]	0
Conv2d-35	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-36	[-1, 512, 16, 16]	1,024
ReLU-37	[-1, 512, 16, 16]	0

Conv2d-38	[-1, 256, 32, 32]	295,168
BatchNorm2d-39	[-1, 256, 32, 32]	512
ReLU-40	[-1, 256, 32, 32]	0
Dropout2d-41	[-1, 256, 32, 32]	0
Conv2d-42	[-1, 256, 32, 32]	590,080
BatchNorm2d-43	[-1, 256, 32, 32]	512
ReLU-44	[-1, 256, 32, 32]	0
ConvTranspose2d-45	[-1, 512, 32, 32]	1,049,088
Conv2d-46	[-1, 256, 32, 32]	1,179,904
BatchNorm2d-47	[-1, 256, 32, 32]	512
ReLU-48	[-1, 256, 32, 32]	0
Dropout2d-49	[-1, 256, 32, 32]	0
Conv2d-50	[-1, 256, 32, 32]	590,080
BatchNorm2d-51	[-1, 256, 32, 32]	512
ReLU-52	[-1, 256, 32, 32]	0
Conv2d-53	[-1, 128, 64, 64]	73,856
BatchNorm2d-54	[-1, 128, 64, 64]	256
ReLU-55	[-1, 128, 64, 64]	0
Dropout2d-56	[-1, 128, 64, 64]	0
Conv2d-57	[-1, 128, 64, 64]	147,584
BatchNorm2d-58	[-1, 128, 64, 64]	256
ReLU-59	[-1, 128, 64, 64]	0
ConvTranspose2d-60	[-1, 512, 64, 64]	1,049,088
Conv2d-61	[-1, 128, 64, 64]	589,952
BatchNorm2d-62	[-1, 128, 64, 64]	256
ReLU-63	[-1, 128, 64, 64]	0
Dropout2d-64	[-1, 128, 64, 64]	0
Conv2d-65	[-1, 128, 64, 64]	147,584
BatchNorm2d-66	[-1, 128, 64, 64]	256
ReLU-67	[-1, 128, 64, 64]	0
Conv2d-68	[-1, 64, 128, 128]	18,496
BatchNorm2d-69	[-1, 64, 128, 128]	128
ReLU-70	[-1, 64, 128, 128]	0
Conv2d-71	[-1, 64, 128, 128]	36,928
BatchNorm2d-72	[-1, 64, 128, 128]	128
ReLU-73	[-1, 64, 128, 128]	0
ConvTranspose2d-74	[-1, 256, 128, 128]	262,400
Conv2d-75	[-1, 64, 128, 128]	147,520
BatchNorm2d-76	[-1, 64, 128, 128]	128
ReLU-77	[-1, 64, 128, 128]	0
Conv2d-78	[-1, 64, 128, 128]	36,928
BatchNorm2d-79	[-1, 64, 128, 128]	128
ReLU-80	[-1, 64, 128, 128]	0
Conv2d-81	[-1, 32, 256, 256]	896
BatchNorm2d-82	[-1, 32, 256, 256]	64
ReLU-83	[-1, 32, 256, 256]	0
Conv2d-84	[-1, 32, 256, 256]	9,248
BatchNorm2d-85	[-1, 32, 256, 256]	64
ReLU-86	[-1, 32, 256, 256]	0
ConvTranspose2d-87	[-1, 128, 256, 256]	65,664
Conv2d-88	[-1, 32, 256, 256]	36,896
BatchNorm2d-89	[-1, 32, 256, 256]	64
ReLU-90	[-1, 32, 256, 256]	0
Conv2d-91	[-1, 32, 256, 256]	9,248
BatchNorm2d-92	[-1, 32, 256, 256]	64
ReLU-93	[-1, 32, 256, 256]	0
Conv2d-94	[-1, 1, 256, 256]	577
BatchNorm2d-95	[-1, 1, 256, 256]	2

=====

Total params: 11,405,859

Trainable params: 11,405,859

Non-trainable params: 0

Input size (MB): 0.75

Forward/backward pass size (MB): 689.50

Params size (MB): 43.51

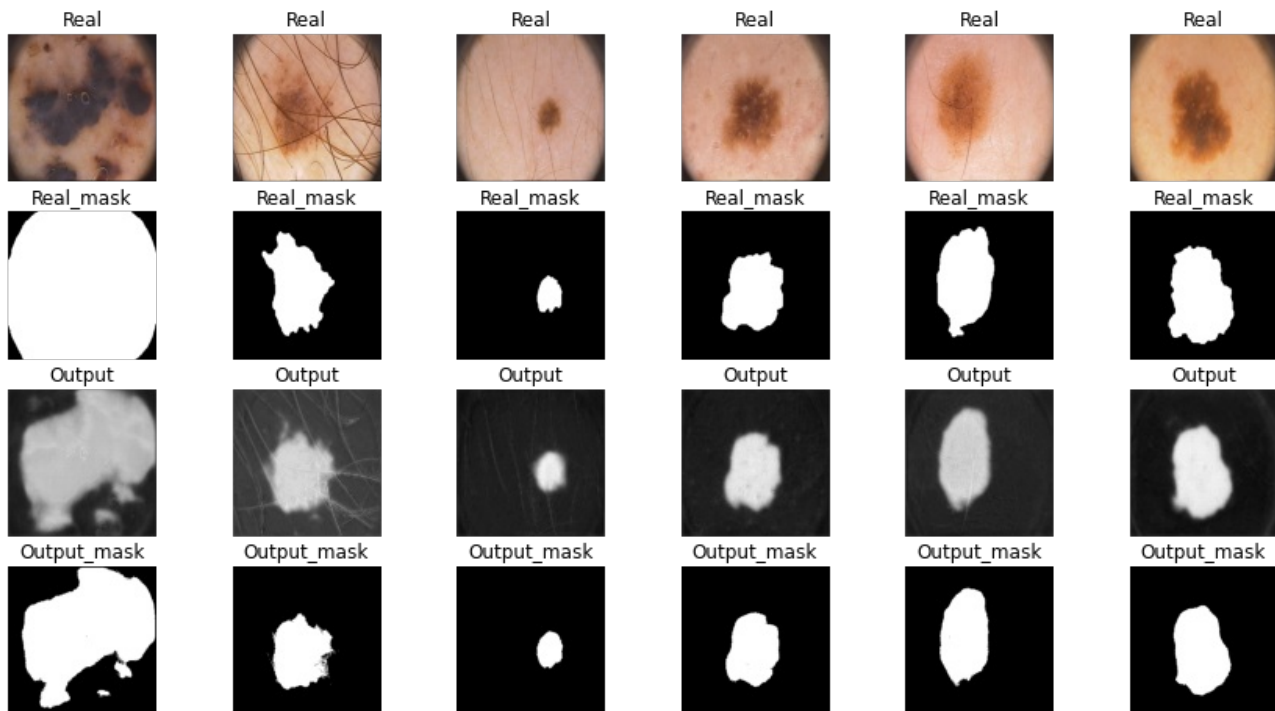
Estimated Total Size (MB): 733.76

In [55]:

```
seed_init()
clear_memory()

UNET2_2_model = UNet2(near_neigh=False).to(device)
opt = optim.Adam(UNET2_2_model.parameters(), lr=lr)
history_UNET2_2 = train('UNET2_2', UNET2_2_model, opt, bce_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.7111, train_score: 0.735, val_loss: 0.7205, val_score: 0.632, max_val_score: 0.684



In [56]:

```
score_model(UNET2_2_model, iou_pytorch, data_val)
```

Out[56]:


0.6840000152587891

In [57]:

```
plot_stat(history_unet_2_2)
```

max score on train: 0.740999968598175 with loss: 0.7117146849632263

max score on val: 0.6840000152587891 with loss: 0.7189147472381592



Сделайте вывод какая из моделей лучше

Dilated convolutions [1 балл]

Еще один из вариантов делать upsampling и downsampling - использовать для этого *dilated convolutions*:

- Yu, Fisher, and Vladlen Koltun. "[Multi-scale context aggregation by dilated convolutions](https://arxiv.org/pdf/1511.07122.pdf)." arXiv preprint arXiv:1511.07122 (2015).

попробуйте написать сеть DilatedUNet, которая использует в одной из предыдущих моделей dilated свертки.

In [0]:

```
class DilatedUNet(nn.Module):
    def __init__(self):
        super().__init__()

        # encoder (downsampling)
        self.enc_conv0 = nn.Sequential(
            nn.Conv2d(3, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
        )

        self.pool0 = nn.Conv2d(32, 32, kernel_size=2, stride=2, padding=1, dilation=2) # 256 -> 128
        self.enc_conv1 = nn.Sequential(
            nn.Conv2d(32, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, kernel_size=3, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
        )

        self.pool1 = nn.Conv2d(64, 64, kernel_size=2, stride=2, padding=1, dilation=2) # 128 -> 64
```

```

self.pool1 = nn.Conv2d(1, 1, kernel_size=1, stride=1, padding=1, dilation=1, groups=1)
self.enc_conv2 = nn.Sequential(
    nn.Conv2d(64, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(128, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
)

self.pool2 = nn.Conv2d(128, 128, kernel_size=2, stride=2, padding=1, dilation=2) # 64 -> 32
self.enc_conv3 = nn.Sequential(
    nn.Conv2d(128, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(256, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
)

self.pool3 = nn.Conv2d(256, 256, kernel_size=2, stride=2, padding=1, dilation=2) # 32 -> 16

# bottleneck
self.bottleneck_conv = nn.Sequential(
    nn.Conv2d(256, 512, kernel_size=3, padding=1),
    nn.BatchNorm2d(512),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(512, 512, kernel_size=3, padding=1),
    nn.BatchNorm2d(512),
    nn.ReLU(inplace=True),
)

# decoder (upsampling)
self.upsample0 = nn.ConvTranspose2d(512, 512, kernel_size=2, stride=2, padding=1, output_padding=1,
dilation=2) # 16 -> 32 # torch.nn.ConvTranspose2d(in_channels=512, out_channels=256, kernel_size=3, stride=2
, padding=1, output_padding=1)
self.dec_conv0 = nn.Sequential(
    nn.Conv2d(512, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(256, 256, kernel_size=3, padding=1),
    nn.BatchNorm2d(256),
    nn.ReLU(inplace=True),
)

# *2 - because concatenation
self.upsample1 = nn.ConvTranspose2d(256*2, 256*2, kernel_size=2, stride=2, padding=1, output_padding
=1, dilation=2) # 32 -> 64
self.dec_conv1 = nn.Sequential(
    nn.Conv2d(256*2, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
    nn.Dropout2d(p=0.3),
    nn.Conv2d(128, 128, kernel_size=3, padding=1),
    nn.BatchNorm2d(128),
    nn.ReLU(inplace=True),
)

# *2 - because concatenation
self.upsample2 = nn.ConvTranspose2d(128*2, 128*2, kernel_size=2, stride=2, padding=1, output_paddin
g=1, dilation=2) # 64 -> 128
self.dec_conv2 = nn.Sequential(
    nn.Conv2d(128*2, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),
    nn.Conv2d(64, 64, kernel_size=3, padding=1),
    nn.BatchNorm2d(64),
    nn.ReLU(inplace=True),
)

# *2 - because concatenation
self.upsample3 = nn.ConvTranspose2d(64*2, 64*2, kernel_size=2, stride=2, padding=1, output_padding=1
, dilation=2) # 128 -> 256
self.dec_conv3 = nn.Sequential(
    nn.Conv2d(64*2, 32, kernel_size=3, padding=1),
    nn.BatchNorm2d(32),

```



```

nn.ReLU(inplace=True),
nn.Conv2d(32, 32, kernel_size=3, padding=1),
nn.BatchNorm2d(32),
nn.ReLU(inplace=True),
)
# *2 - because concatenation
self.dec_final = nn.Sequential(nn.Conv2d(32*2, 1, kernel_size=3, padding=1),
                               nn.BatchNorm2d(1)
                              )

def forward(self, x):
    # encoder
    e0 = self.pool0(self.enc_conv0(x))
    e1 = self.pool1(self.enc_conv1(e0))
    e2 = self.pool2(self.enc_conv2(e1))
    e3 = self.pool3(self.enc_conv3(e2))

    # bottleneck
    b = self.bottleneck_conv(e3)

    # decoder
    # print('enc_part: ', self.enc_conv3(e2).shape, '+', 'dec_part: ', self.dec_conv0(self.upsample0(b)).
shape)
    d0 = torch.cat((self.enc_conv3(e2), self.dec_conv0(self.upsample0(b))), 1)
    # print('enc_part: ', self.enc_conv2(e1).shape, '+', 'dec_part: ', self.dec_conv1(self.upsample1(d0))
.shape)
    d1 = torch.cat((self.enc_conv2(e1), self.dec_conv1(self.upsample1(d0))), 1)
    # print('enc_part: ', self.enc_conv1(e0).shape, '+', 'dec_part: ', self.dec_conv2(self.upsample2(d1))
.shape)
    d2 = torch.cat((self.enc_conv1(e0), self.dec_conv2(self.upsample2(d1))), 1)
    # print('enc_part: ', self.enc_conv0(x).shape, '+', 'dec_part: ', self.dec_conv3(self.upsample3(d2)).
shape)
    d3 = torch.cat((self.enc_conv0(x), self.dec_conv3(self.upsample3(d2))), 1)
    # print(d3.shape)
    dfin = self.dec_final(d3)    # no activation
    # print(dfin.shape)
    return dfin

```

In [59]:

```
summary(DilatedUNet().to(device), input_size=(3, 256, 256))
```

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 32, 256, 256]	896
BatchNorm2d-2	[-1, 32, 256, 256]	64
ReLU-3	[-1, 32, 256, 256]	0
Conv2d-4	[-1, 32, 256, 256]	9,248
BatchNorm2d-5	[-1, 32, 256, 256]	64
ReLU-6	[-1, 32, 256, 256]	0
Conv2d-7	[-1, 32, 128, 128]	4,128
Conv2d-8	[-1, 64, 128, 128]	18,496
BatchNorm2d-9	[-1, 64, 128, 128]	128
ReLU-10	[-1, 64, 128, 128]	0
Conv2d-11	[-1, 64, 128, 128]	36,928
BatchNorm2d-12	[-1, 64, 128, 128]	128
ReLU-13	[-1, 64, 128, 128]	0
Conv2d-14	[-1, 64, 64, 64]	16,448
Conv2d-15	[-1, 128, 64, 64]	73,856
BatchNorm2d-16	[-1, 128, 64, 64]	256
ReLU-17	[-1, 128, 64, 64]	0
Dropout2d-18	[-1, 128, 64, 64]	0
Conv2d-19	[-1, 128, 64, 64]	147,584
BatchNorm2d-20	[-1, 128, 64, 64]	256
ReLU-21	[-1, 128, 64, 64]	0
Conv2d-22	[-1, 128, 32, 32]	65,664
Conv2d-23	[-1, 256, 32, 32]	295,168
BatchNorm2d-24	[-1, 256, 32, 32]	512
ReLU-25	[-1, 256, 32, 32]	0
Dropout2d-26	[-1, 256, 32, 32]	0
Conv2d-27	[-1, 256, 32, 32]	590,080
BatchNorm2d-28	[-1, 256, 32, 32]	512
ReLU-29	[-1, 256, 32, 32]	0
Conv2d-30	[-1, 256, 16, 16]	262,400
Conv2d-31	[-1, 512, 16, 16]	1,180,160
BatchNorm2d-32	[-1, 512, 16, 16]	1,024
ReLU-33	[-1, 512, 16, 16]	0

Dropout2d-34	[-1, 512, 16, 16]	0
Conv2d-35	[-1, 512, 16, 16]	2,359,808
BatchNorm2d-36	[-1, 512, 16, 16]	1,024
ReLU-37	[-1, 512, 16, 16]	0
Conv2d-38	[-1, 256, 32, 32]	295,168
BatchNorm2d-39	[-1, 256, 32, 32]	512
ReLU-40	[-1, 256, 32, 32]	0
Dropout2d-41	[-1, 256, 32, 32]	0
Conv2d-42	[-1, 256, 32, 32]	590,080
BatchNorm2d-43	[-1, 256, 32, 32]	512
ReLU-44	[-1, 256, 32, 32]	0
ConvTranspose2d-45	[-1, 512, 32, 32]	1,049,088
Conv2d-46	[-1, 256, 32, 32]	1,179,904
BatchNorm2d-47	[-1, 256, 32, 32]	512
ReLU-48	[-1, 256, 32, 32]	0
Dropout2d-49	[-1, 256, 32, 32]	0
Conv2d-50	[-1, 256, 32, 32]	590,080
BatchNorm2d-51	[-1, 256, 32, 32]	512
ReLU-52	[-1, 256, 32, 32]	0
Conv2d-53	[-1, 128, 64, 64]	73,856
BatchNorm2d-54	[-1, 128, 64, 64]	256
ReLU-55	[-1, 128, 64, 64]	0
Dropout2d-56	[-1, 128, 64, 64]	0
Conv2d-57	[-1, 128, 64, 64]	147,584
BatchNorm2d-58	[-1, 128, 64, 64]	256
ReLU-59	[-1, 128, 64, 64]	0
ConvTranspose2d-60	[-1, 512, 64, 64]	1,049,088
Conv2d-61	[-1, 128, 64, 64]	589,952
BatchNorm2d-62	[-1, 128, 64, 64]	256
ReLU-63	[-1, 128, 64, 64]	0
Dropout2d-64	[-1, 128, 64, 64]	0
Conv2d-65	[-1, 128, 64, 64]	147,584
BatchNorm2d-66	[-1, 128, 64, 64]	256
ReLU-67	[-1, 128, 64, 64]	0
Conv2d-68	[-1, 64, 128, 128]	18,496
BatchNorm2d-69	[-1, 64, 128, 128]	128
ReLU-70	[-1, 64, 128, 128]	0
Conv2d-71	[-1, 64, 128, 128]	36,928
BatchNorm2d-72	[-1, 64, 128, 128]	128
ReLU-73	[-1, 64, 128, 128]	0
ConvTranspose2d-74	[-1, 256, 128, 128]	262,400
Conv2d-75	[-1, 64, 128, 128]	147,520
BatchNorm2d-76	[-1, 64, 128, 128]	128
ReLU-77	[-1, 64, 128, 128]	0
Conv2d-78	[-1, 64, 128, 128]	36,928
BatchNorm2d-79	[-1, 64, 128, 128]	128
ReLU-80	[-1, 64, 128, 128]	0
Conv2d-81	[-1, 32, 256, 256]	896
BatchNorm2d-82	[-1, 32, 256, 256]	64
ReLU-83	[-1, 32, 256, 256]	0
Conv2d-84	[-1, 32, 256, 256]	9,248
BatchNorm2d-85	[-1, 32, 256, 256]	64
ReLU-86	[-1, 32, 256, 256]	0
ConvTranspose2d-87	[-1, 128, 256, 256]	65,664
Conv2d-88	[-1, 32, 256, 256]	36,896
BatchNorm2d-89	[-1, 32, 256, 256]	64
ReLU-90	[-1, 32, 256, 256]	0
Conv2d-91	[-1, 32, 256, 256]	9,248
BatchNorm2d-92	[-1, 32, 256, 256]	64
ReLU-93	[-1, 32, 256, 256]	0
Conv2d-94	[-1, 1, 256, 256]	577
BatchNorm2d-95	[-1, 1, 256, 256]	2

```

=====
Total params: 11,405,859
Trainable params: 11,405,859
Non-trainable params: 0

```

```

-----
Input size (MB): 0.75
Forward/backward pass size (MB): 689.50
Params size (MB): 43.51
Estimated Total Size (MB): 733.76
-----

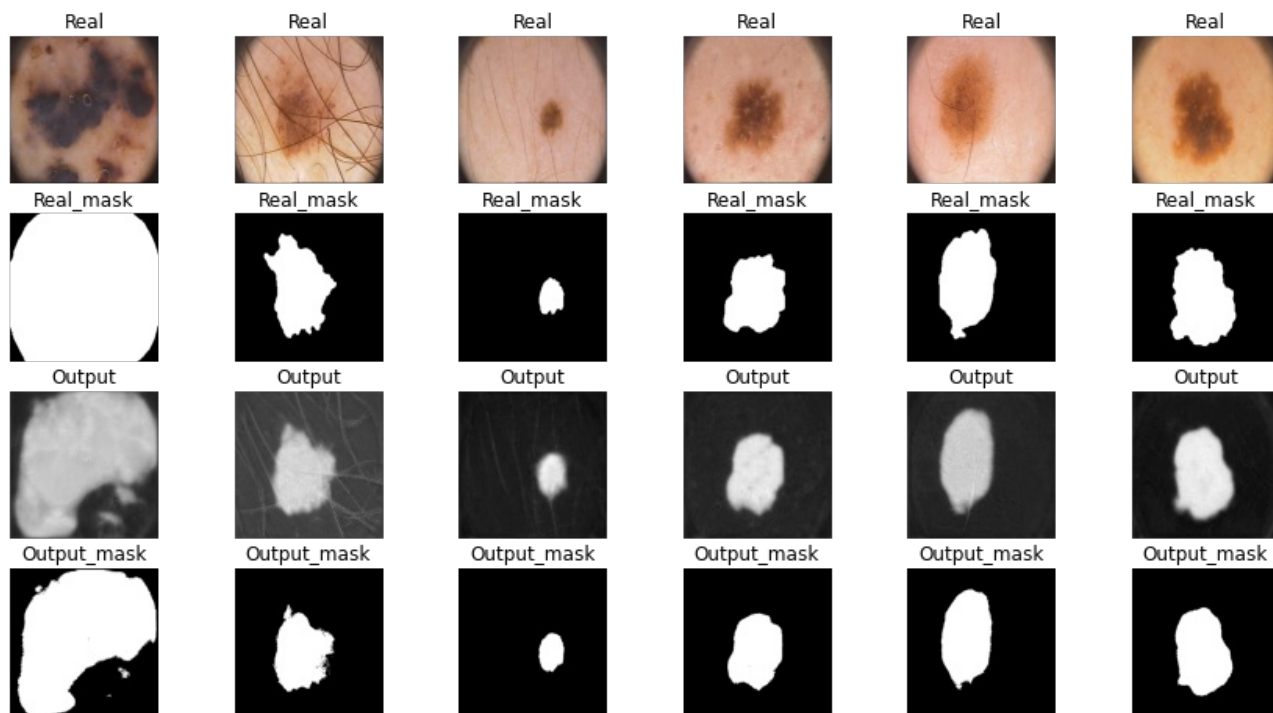
```

In [60]:

```
seed_init()
clear_memory()

unetdil_model = DilatedUNet().to(device)
opt = optim.Adam(unetdil_model.parameters(), lr=lr)
history_unetdil = train('unet_dil', unetdil_model, opt, bce_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.7109, train_score: 0.743, val_loss: 0.7206, val_score: 0.644, max_val_score: 0.664





In [61]:

```
plot_stat(history_unetdil)
```

max score on train: 0.7490000277757645 with loss: 0.7107290476560593

max score on val: 0.6639999747276306 with loss: 0.7213621139526367



In [62]:

```
score_model(unetdil_model, iou_pytorch, data_val)
```

Out[62]:

0.664000004529953

Отчет (5 баллов):

Ниже предлагается написать отчет о проделанной работе и построить графики для потерь, метрик на валидации и тесте.

Аккуратно сравните модели между собой и соберите наилучшую архитектуру. Проверьте каждую модель с различными потерями. Мы не ограничиваем вас в формате отчета, но проверяющий должен отчетливо понять для чего построен каждый график, какие выводы вы из него сделали и какой общий вывод можно сделать на основании данных моделей. Если вы захотите добавить что-то еще, чтобы увеличить шансы получения максимального балла, то добавляйте отдельное сравнение.

In [63]:

```
# Loss and Score on test
```

```
stat_on_test = dict()
losses_fn = [bce_loss, dice_loss, focal_loss, tversky_loss(beta=0.3), bce_loss, bce_loss, bce_loss, bce_loss]
names = ['bce', 'dice', 'focal', 'tversky', 'maxpool(2,2) & upsamp(bilin)', 'maxpool(3,2,1) & upsamp(nearest)', 'conv(2,2) & conv trans(2, 2)', 'dilation']
models = [model_bce, model_dice, model_focal, model_tversky, unet_model, unet2_1_model, unet2_2_model, unetdil_model]

for name, model, loss_fn in zip(names, models, losses_fn):
    stat_on_test[name] = dict(zip(['loss', 'score'], eval(model, data_ts, loss_fn, iou_pytorch, threshold=.7)))
```

```
/usr/local/lib/python3.6/dist-packages/torch/nn/functional.py:2973: UserWarning: Default upsampling behavior when mode=bilinear is changed to align_corners=False since 0.4.0. Please specify align_corners=True if the old behavior is desired. See the documentation of nn.Upsample for details.
```

```
"See the documentation of nn.Upsample for details.".format(mode))
```

In [64]:

```
plt.figure(figsize=(15, 10))
i = 1
for name, history in zip(['bce', 'dice', 'focal', 'tversky'], [history_bce, history_dice, history_focal, history_tversky]):
    train_loss, train_score, val_loss, val_score = zip(*history)
    max_score_train = np.max(train_score)
    max_score_loss_train = train_loss[np.argmax(train_score)]
    max_score_val = np.max(val_score)
    max_score_loss_val = val_loss[np.argmax(val_score)]
    ts_score, ts_loss = stat_on_test[name]['score'], stat_on_test[name]['loss']

    print(f'\nloss_fn {name}: \nmax score on train, eph = {np.argmax(train_score)}: {np.round(max_score_train, 4)} with loss: {np.round(max_score_loss_train, 4)}')
    print(f'mean 5 last ephs score on train: {np.round(np.mean(train_score[max_epochs-5:]), 4)} with loss: {np.round(np.mean(train_loss[max_epochs-5:]), 4)}')
    print(f'\nmax score on val, eph = {np.argmax(val_score)}: {np.round(max_score_val, 4)} with loss: {np.round(max_score_loss_val, 4)}')
    print(f'mean 5 last ephs score on val: {np.round(np.mean(val_score[max_epochs-5:]), 4)} with loss: {np.round(np.mean(val_loss[max_epochs-5:]), 4)}')
    print(f'\nscore on test: {np.round(ts_score, 4)} with loss: {np.round(ts_loss, 4)}')
    print('='*100)

    plt.subplot(4, 2, i)
    plt.plot(train_score, label=f"train_score_{name}")
    plt.plot(val_score, label=f"val_score_{name}")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("score")

    plt.subplot(4, 2, i+1)
    plt.plot(train_loss, label=f"train_loss_{name}")
    plt.plot(val_loss, label=f"val_loss_{name}")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("loss")
    i += 2

plt.show()
```

loss_fn bce:
max score on train, eph = 39: 0.667 with loss: 0.7153
mean 5 last ephs score on train: 0.6512 with loss: 0.7148

max score on val, eph = 29: 0.668 with loss: 0.7214
mean 5 last ephs score on val: 0.658 with loss: 0.7181

score on test: 0.6618 with loss: 0.7338

loss_fn dice:
max score on train, eph = 38: 0.645 with loss: 0.7093
mean 5 last ephs score on train: 0.6184 with loss: 0.7096

max score on val, eph = 29: 0.642 with loss: 0.7307
mean 5 last ephs score on val: 0.6024 with loss: 0.7353

score on test: 0.6061 with loss: 0.7365

loss_fn focal:
max score on train, eph = 36: 0.647 with loss: 0.2153
mean 5 last ephs score on train: 0.6358 with loss: 0.2155

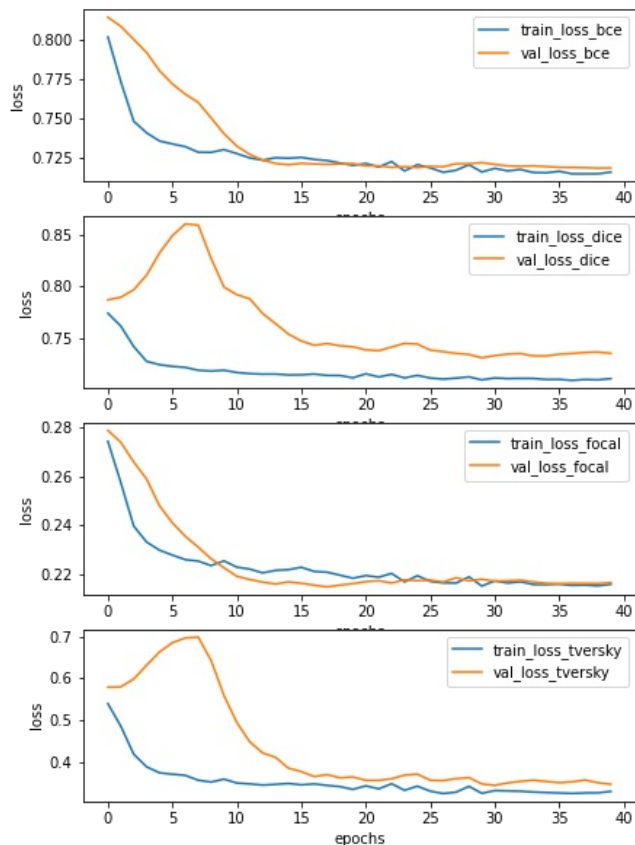
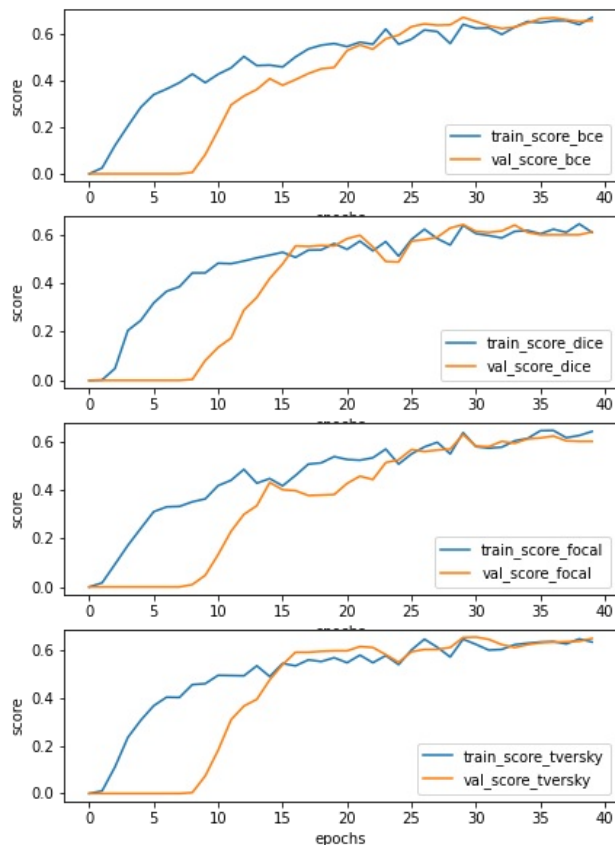
max score on val, eph = 29: 0.63 with loss: 0.2178
mean 5 last ephs score on val: 0.6096 with loss: 0.2161

score on test: 0.6655 with loss: 0.2242

loss_fn tversky:
max score on train, eph = 38: 0.65 with loss: 0.3271
mean 5 last ephs score on train: 0.6384 with loss: 0.3274

max score on val, eph = 30: 0.658 with loss: 0.3448
mean 5 last ephs score on val: 0.64 with loss: 0.3522

score on test: 0.653 with loss: 0.3332



In [65]:

```
plt.figure(figsize=(15, 10))
i = 1
for name, history in zip(['maxpool(2,2) & upsamp(bilin)', 'maxpool(3,2,1) & upsamp(nearest)', 'conv(2,2) & c
onv trans(2, 2)', 'dilation'], [history_unet, history_unet_2_1, history_unet_2_2, history_unetdil]):
    train_loss, train_score, val_loss, val_score = zip(*history)
    max_score_train = np.max(train_score)
    max_score_loss_train = train_loss[np.argmax(train_score)]
    max_score_val = np.max(val_score)
    max_score_loss_val = val_loss[np.argmax(val_score)]
    ts_score, ts_loss = stat_on_test[name]['score'], stat_on_test[name]['loss']

    print(f'\nunet with {name}: \nmax score on train, eph = {np.argmax(train_score)}: {np.round(max_score_tr
ain, 4)} with loss: {np.round(max_score_loss_train, 4)}')
    print(f'mean 5 last ephs score on train: {np.round(np.mean(train_score[max_epochs-5:]), 4)} with loss: {
np.round(np.mean(train_loss[max_epochs-5:]), 4)}')
    print(f'\nmax score on val, eph = {np.argmax(val_score)}: {np.round(max_score_val, 4)} with loss: {np.ro
und(max_score_loss_val, 4)}')
    print(f'mean 5 last ephs score on val: {np.round(np.mean(val_score[max_epochs-5:]), 4)} with loss: {np.r
ound(np.mean(val_loss[max_epochs-5:]), 4)}')
    print(f'\nscore on test: {np.round(ts_score, 4)} with loss: {np.round(ts_loss, 4)}')
    print('='*100)

    plt.subplot(4, 2, i)
    plt.plot(train_score, label=f"train_score_{name}")
    plt.plot(val_score, label=f"val_score_{name}")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("score")

    plt.subplot(4, 2, i+1)
    plt.plot(train_loss, label=f"train_loss_{name}")
    plt.plot(val_loss, label=f"val_loss_{name}")
    plt.legend(loc='best')
    plt.xlabel("epochs")
    plt.ylabel("loss")
    i+=2

plt.show()
```

unet with maxpool(2,2) & upsamp(bilin):
max score on train, eph = 38: 0.744 with loss: 0.7106
mean 5 last ephs score on train: 0.714 with loss: 0.7123

max score on val, eph = 29: 0.704 with loss: 0.7212
mean 5 last ephs score on val: 0.6872 with loss: 0.7179

score on test: 0.6872 with loss: 0.7301

=====

unet with maxpool(3,2,1) & upsamp(nearest):
max score on train, eph = 33: 0.758 with loss: 0.7112
mean 5 last ephs score on train: 0.731 with loss: 0.7115

max score on val, eph = 28: 0.718 with loss: 0.7201
mean 5 last ephs score on val: 0.6908 with loss: 0.7171

score on test: 0.6963 with loss: 0.7318

=====

unet with conv(2,2) & conv trans(2, 2):
max score on train, eph = 35: 0.741 with loss: 0.7117
mean 5 last ephs score on train: 0.7244 with loss: 0.7118

max score on val, eph = 30: 0.684 with loss: 0.7189
mean 5 last ephs score on val: 0.6428 with loss: 0.7193

score on test: 0.6963 with loss: 0.725

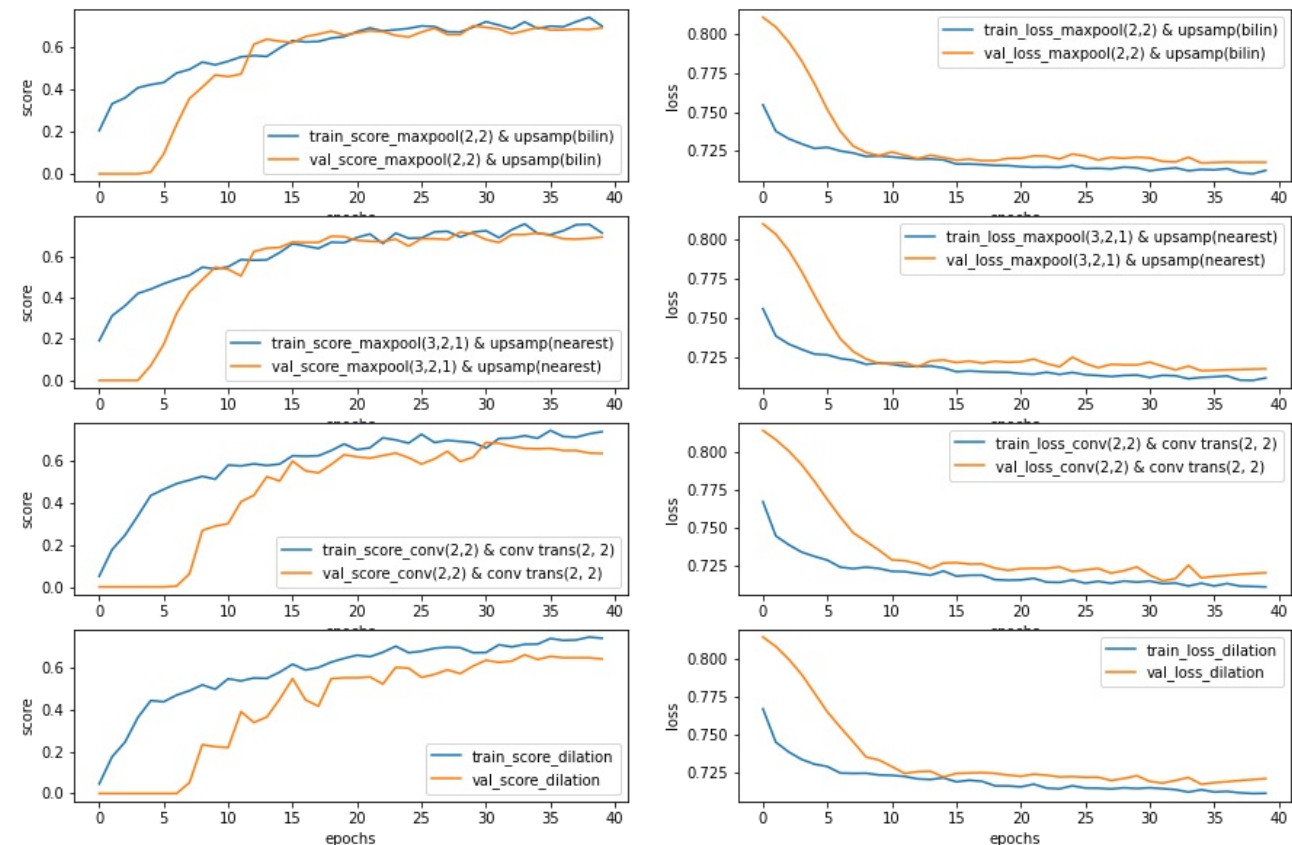
=====

unet with dilation:
max score on train, eph = 38: 0.749 with loss: 0.7107
mean 5 last ephs score on train: 0.7404 with loss: 0.7113

max score on val, eph = 33: 0.664 with loss: 0.7214
mean 5 last ephs score on val: 0.65 with loss: 0.7193

score on test: 0.6968 with loss: 0.7248

=====



Выводы

Segnet with different Losses

	BCE loss	dice loss	focal loss	tversky loss (beta=0.7)
max val score	0.668	0.642	0.63	0.658
mean 5 last ephs val score	0.658	0.6024	0.6908	0.64
test score	0.6618	0.6061	0.6655	0.653

В таблице приведена метрика IoU на валидации при обучении, а также на тесте после обучения:
model = SegNet
optimaizer = Adam
lr = 1e-4
scheduler = gamma = 0.1, step_size = 35
epochs = 40

Заключения по таблице:
Лучший результат наблюдается при использовании целевой функции **focal loss**:

	focal loss
max val score	0.63
mean 5 last ephs val score	0.6908
test score	0.6655

Unet with Downsampling & Upsampling

	downsample & upsampling	downsample & upsampling	downsample & upsampling	Dilated convolutions downsample & upsampling
Unet with	MaxPool2d(kernel_size = 2, stride = 2) Upsample(scale_factor = 2, mode = "bilinear")	MaxPool2d(kernel_size = 3, stride = 2, padding = 1) Upsample(scale_factor = 2, mode = 'nearest')	Conv2d(input, output, kernel_size = 2, stride = 2) ConvTranspose2d(input, output, kernel_size = 2, stride = 2)	Conv2d(input, output, kernel_size = 2, stride = 2, padding = 1, dilation = 2) ConvTranspose2d(input, output, kernel_size = 2, stride = 2, padding = 1, output_padding = 1, dilation = 2)
max val score	0.704	0.718	0.684	0.664
mean 5 last ephs val score	0.6872	0.6908	0.6428	0.65
test score	0.6872	0.6963	0.6963	0.6968

В таблице приведена метрика IoU на валидации при обучении, а также на тесте после обучения:
model = Unet
optimaizer = Adam
scheduler = gamma = 0.1, step_size = 35
lr = 1e-4
epochs = 40

Заклучения по таблице:
Лучший результат наблюдается при использовании в качестве
Downsampling - Conv2d(input, output, kernel_size = 2, stride = 2, padding = 1, dilation = 2),
Upsampling - ConvTranspose2d(input, output, kernel_size = 2, stride = 2, padding = 1, output_padding = 1, dilation = 2):

	Dilated convolutions downsample & upsampling
Unet with	Conv2d(input, output, kernel_size = 2, stride = 2, padding = 1, dilation = 2) ConvTranspose2d(input, output, kernel_size = 2, stride = 2, padding = 1, output_padding = 1, dilation = 2)
max val score	0.664
mean 5 last ephs val score	0.65
test score	0.6968

Результаты на тесте имеет наивысший приоритет:
test score = 0.6968

Наилучшая архитектура

Unet with Dilated convolutions

Downsampling - `Conv2d(input, output, kernel_size = 2, stride = 2, padding = 1, dilation = 2),`

Upsampling - `ConvTranspose2d(input, output, kernel_size = 2, stride = 2, padding = 1, output_padding = 1, dilation = 2)`

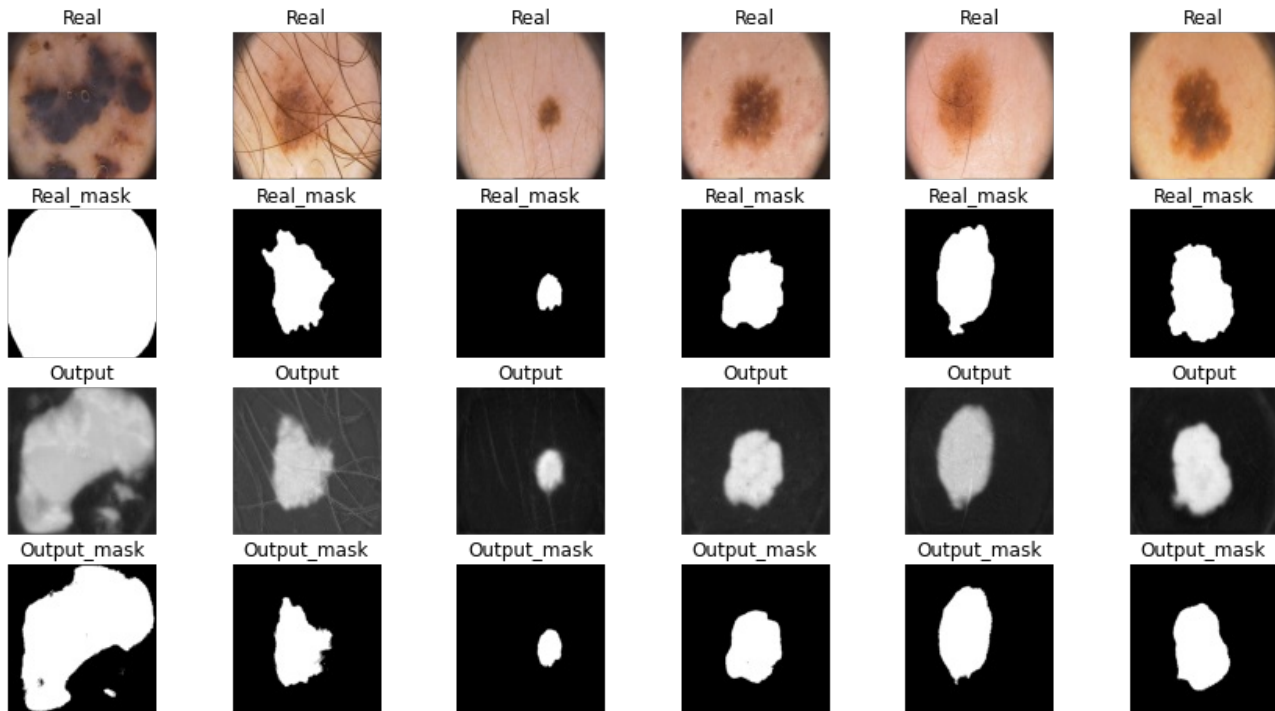
Loss - *focal loss*

In [66]:

```
seed_init()
clear_memory()

model_final = DilatedUNet().to(device)
op = optim.Adam(model_final.parameters(), lr=lr)
history_final = train('final', model_final, op, focal_loss, iou_pytorch, max_epochs, data_tr, data_val)
```

40 / 40 - train_loss: 0.2129, train_score: 0.726, val_loss: 0.2159, val_score: 0.608, max_val_score: 0.69




In [67]:

```
plot_stat(history_final)
```

max score on train: 0.7260000109672546 with loss: 0.21290815249085426

max score on val: 0.6899999976158142 with loss: 0.2169673666357994



In [70]:

```
print('the best score on val: ')\nscore_model(model_final, iou_pytorch, data_val)
```

the best score on val:

Out[70]:

0.6899999678134918

In [69]:

```
train_loss_final, train_score_final, val_loss_final, val_score_final = zip(*history_final)\nmodel_final_ts_score, model_final_ts_loss = eval(model_final, data_ts, bce_loss, iou_pytorch, threshold=.8)\n\nprint(f'\\nfinalmodel: \\nmax score on train, eph = {np.argmax(train_score_final): {np.round(np.max(train_score_final), 4)} with loss: {np.round(train_loss_final[np.argmax(train_score_final)], 4)}')\nprint(f'mean 5 last ephs score on train: {np.round(np.mean(train_score_final[max_epochs-5:]), 4)} with loss: {np.round(np.mean(train_loss_final[max_epochs-5:]), 4)}')\nprint(f'\\nmax score on val, eph = {np.argmax(val_score_final): {np.round(np.max(val_score_final), 4)} with loss: {np.round(val_loss_final[np.argmax(val_score_final)], 4)}')\nprint(f'mean 5 last ephs score on val: {np.round(np.mean(val_score_final[max_epochs-5:]), 4)} with loss: {np.round(np.mean(val_loss_final[max_epochs-5:]), 4)}')\nprint(f'\\ntest score: {model_final_ts_score}, test loss: {model_final_ts_loss}')
```

finalmodel:

max score on train, eph = 38: 0.726 with loss: 0.2129

mean 5 last ephs score on train: 0.716 with loss: 0.2132

max score on val, eph = 33: 0.69 with loss: 0.217

mean 5 last ephs score on val: 0.6212 with loss: 0.2157

test score: 0.7242080569267273, test loss: 0.5523333251476288

Такая архитектура даёт test score = 0.7242